Exam Questions CKA

Certified Kubernetes Administrator (CKA) Program

https://www.2passeasy.com/dumps/CKA/

NEW QUESTION 1
Create a pod with environment variables as var1=value1.Check the environment variable in pod

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl run nginx --image=nginx --restart=Never --env=var1=value1
then
kubectl exec -it nginx -- env
or
kubectl exec -it nginx -- sh -c 'echo $var1'
or
kubectl describe po nginx | grep value1

NEW QUESTION 2
Create a deployment as follows:
- Name:nginx-random
- Exposed via a servicenginx-random
- Ensure that the service & podare accessible via theirrespective DNS records
- The container(s) within anypod(s) running as a part of thisdeployment should use thenginxImage

Next, use the utilitynslookupto lookup the DNS records of the service &pod and write the output to /opt/KUNW00601/service.dnsand/opt/KUNW00601/pod.dnsrespectively.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Solution:
F:\Work\Data Entry Work\Data Entry\20200827\CKA\17 C.JPG

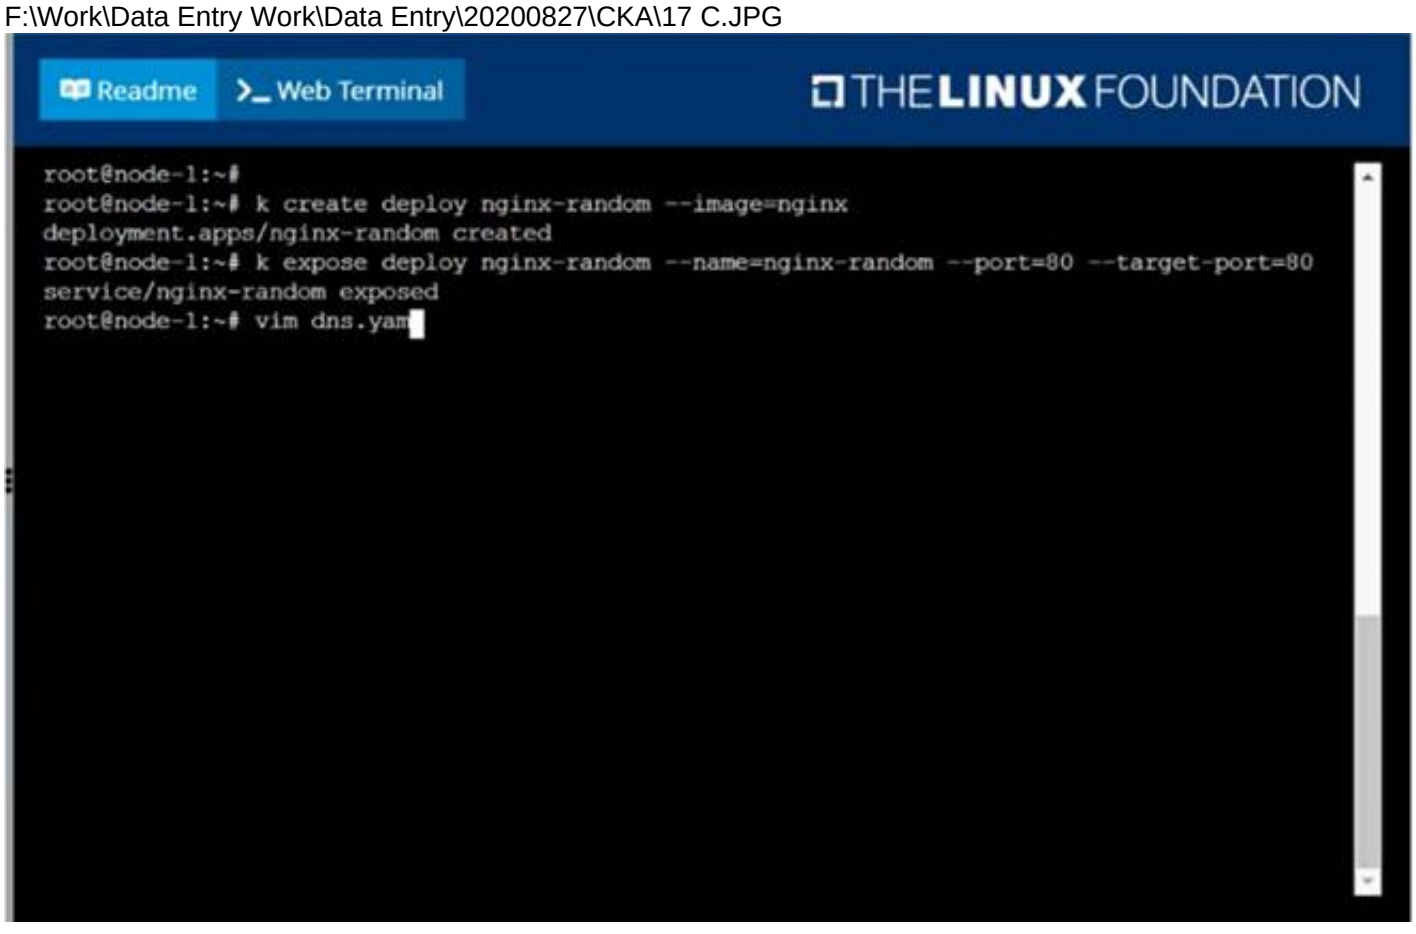

F:\Work\Data Entry Work\Data Entry\20200827\CKA\17 D.JPG

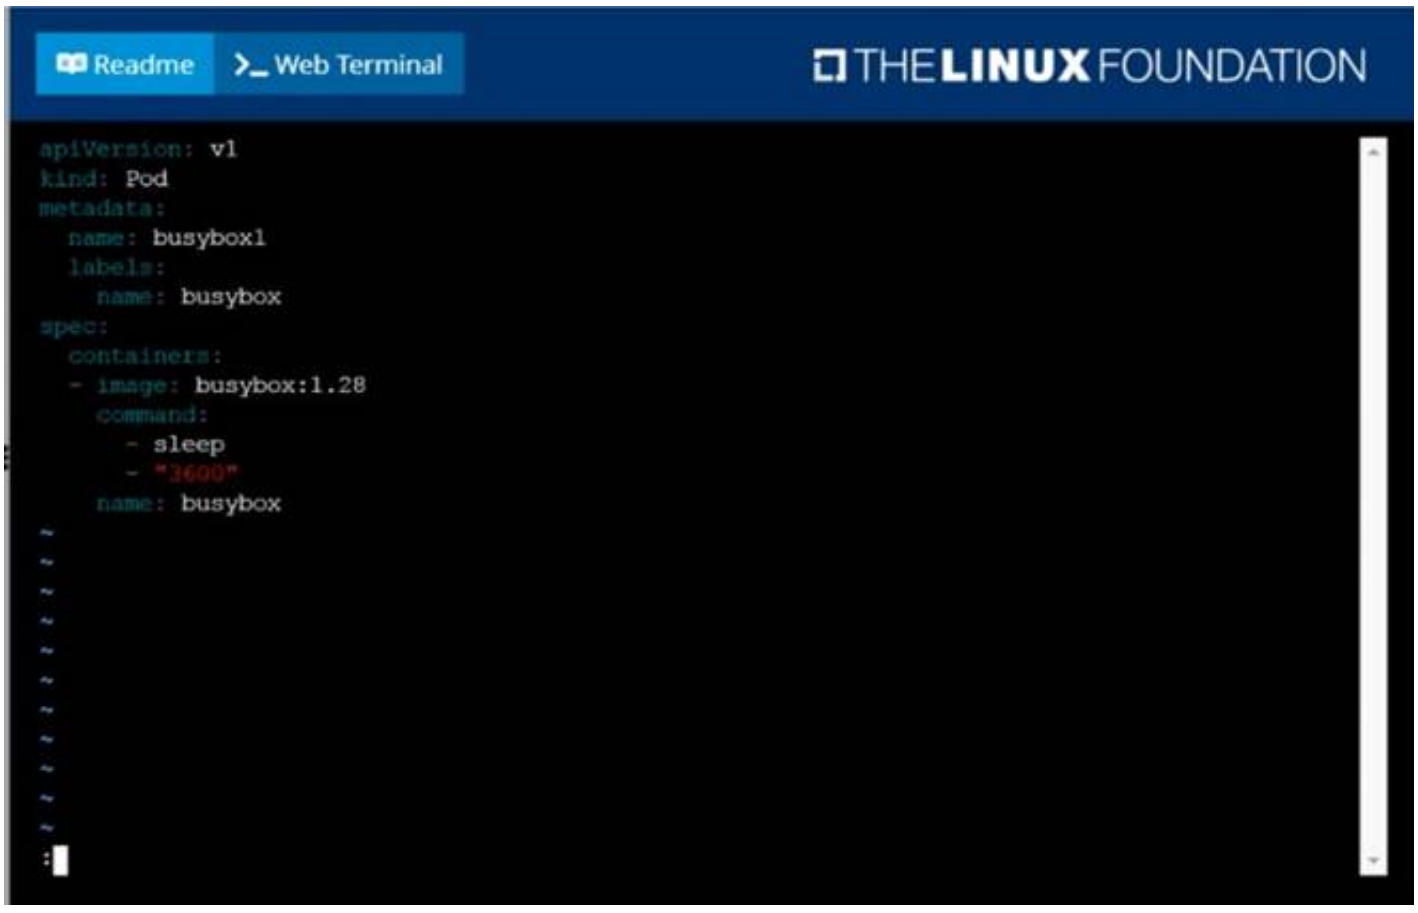

F:\Work\Data Entry Work\Data Entry\20200827\CKA\17 E.JPG

NEW QUESTION 3
Given a partially-functioningKubernetes cluster, identifysymptoms of failure on the cluster.
Determine the node, the failingservice, and take actions to bring upthe failed service and restore thehealth of the cluster. Ensure that anychanges are made permanently.
You cansshto the relevant Inodes (bk8s-master-0orbk8s-node-0) using:
[student@node-1] $ ssh<nodename>
You can assume elevatedprivileges on any node in thecluster with the followingcommand:
[student@nodename] $ | sudo ?Ci

A. Mastered
B. Not Mastered

Answer: A

Explanation:

solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\23 C.JPG

Readme | Web Terminal — THE LINUX FOUNDATION

```
root@node-1:~#
root@node-1:~# kubectl config use-context bk8s
Switched to context "bk8s".
root@node-1:~# ssh bk8s-master-0
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.4.0-1109-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 * Are you ready for Kubernetes 1.19? It's nearly here! Try RC3 with
   sudo snap install microk8s --channel=1.19/candidate --classic

   https://microk8s.io/ has docs and details.

4 packages can be updated.
1 update is a security update.

New release '18.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.


student@bk8s-master-0:~$ sudo -i
root@bk8s-master-0:~# vim /var/lib/kubelet/config.yaml
```

F:\Work\Data Entry Work\Data Entry\20200827\CKA\23 D.JPG

Readme | Web Terminal — THE LINUX FOUNDATION

```
authorization:
  mode: Webhook
  webhook:
    cacheAuthorizedTTL: 0s
    cacheUnauthorizedTTL: 0s
clusterDNS:
- 10.96.0.10
clusterDomain: cluster.local
cpuManagerReconcilePeriod: 0s
evictionPressureTransitionPeriod: 0s
fileCheckFrequency: 0s
healthzBindAddress: 127.0.0.1
healthzPort: 10248
httpCheckFrequency: 0s
imageMinimumGCAge: 0s
kind: KubeletConfiguration
nodeStatusReportFrequency: 0s
nodeStatusUpdateFrequency: 0s
rotateCertificates: true
runtimeRequestTimeout: 0s
staticPodPath: /etc/kubernetes/manifests
streamingConnectionIdleTimeout: 0s
syncFrequency: 0s
volumeStatsAggPeriod: 0s
:wq
```

F:\Work\Data Entry Work\Data Entry\20200827\CKA\23 E.JPG

Readme | Web Terminal — THE LINUX FOUNDATION

```
   https://microk8s.io/ has docs and details.

4 packages can be updated.
1 update is a security update.

New release '18.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.


student@bk8s-master-0:~$ sudo -i
root@bk8s-master-0:~# vim /var/lib/kubelet/config.yaml
root@bk8s-master-0:~# systemctl  restart kubelet
root@bk8s-master-0:~# systemctl  enable kubelet
root@bk8s-master-0:~# kubectl  get nodes

NAME           STATUS   ROLES    AGE   VERSION
bk8s-master-0  Ready    master   77d   v1.18.2
bk8s-node-0    Ready    <none>   77d   v1.18.2
root@bk8s-master-0:~#
root@bk8s-master-0:~# exit
logout
student@bk8s-master-0:~$ exit
logout
Connection to 10.250.4.77 closed.
root@node-1:~#
```

NEW QUESTION 4
Create a pod as follows:
- Name:mongo
- Using Image:mongo
- In anew Kubernetes namespacenamed:my-website

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\9 B.JPG

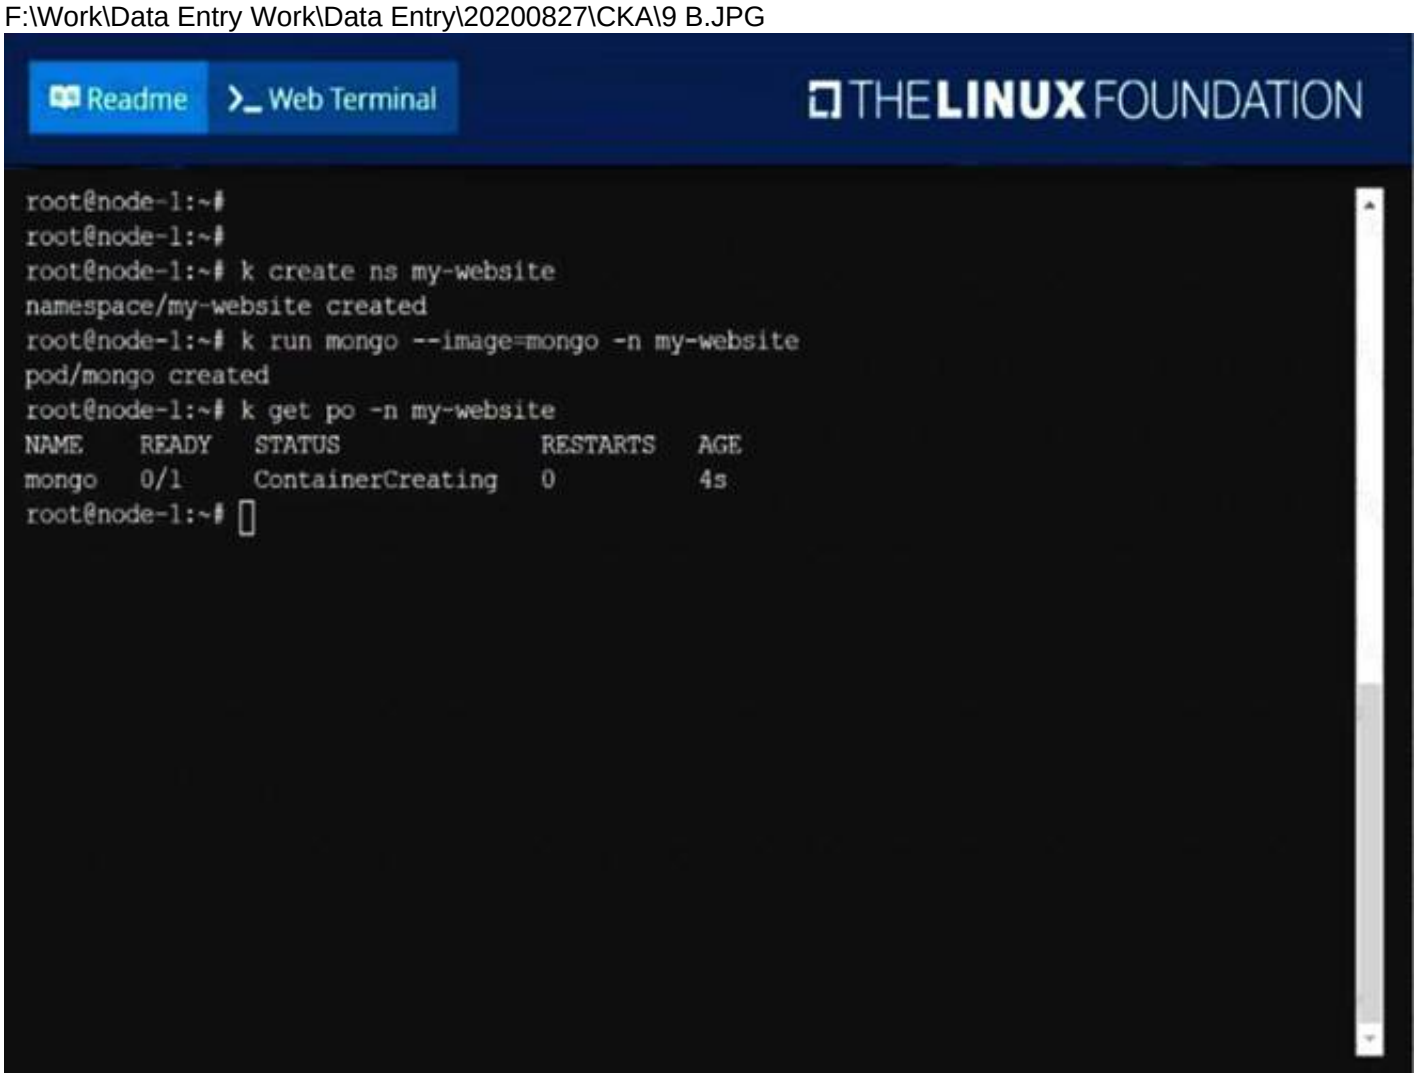

NEW QUESTION 5
List ??nginx-dev?? and ??nginx-prod?? pod and delete those pods

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubect1 get pods -o wide
kubectl delete po ??nginx-dev??kubectl delete po ??nginx-prod??

NEW QUESTION 6
Create a pod with image nginx called nginx and allow traffic on port 80

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectlrun nginx --image=nginx --restart=Never --port=80

NEW QUESTION 7
Create a nginx pod with label env=test in engineering namespace

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-run -o yaml > nginx-pod.yaml

kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-run -o yaml | kubectl create -nengineering-f ?C
YAML File: apiVersion: v1 kind: Pod metadata: name: nginx
namespace: engineering labels:
env: test spec: containers:
- name: nginx image: nginx
imagePullPolicy: IfNotPresent restartPolicy: Never
kubectl create -f nginx-pod.yaml

NEW QUESTION 8
Create a persistent volume with nameapp-data, of capacity2Giandaccess modeReadWriteMany. Thetype of volume ishostPathand itslocation is/srv/app-data.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
Persistent Volume
A persistent volume is a piece of storage in aKubernetes cluster. PersistentVolumes are a cluster-level resource like nodes, which don??t belong to any namespace. It is provisioned by the administrator and has a particular file size. This way, a developer deploying their app on Kubernetes need not knowthe underlying infrastructure. When the developer needs a certain amount of persistent storage for their application, the system administrator configures the cluster so that they consume the PersistentVolume provisioned in an easy way.
Creating PersistentVolume
kind: PersistentVolumeapiVersion: v1metadata:name:app-dataspec:capacity: # defines the capacity of PV we are creatingstorage:2Gi#the amount of storage we are tying to claimaccessModes: # defines the rights of the volumewe are creating-ReadWriteManyhostPath:path: "/srv/app-data" # path to which we are creating the volume
Challenge
> Create a Persistent Volume namedapp-data, with access modeReadWriteMany, storage classname shared,2Giof storage capacity and the host path/srv/app-data.

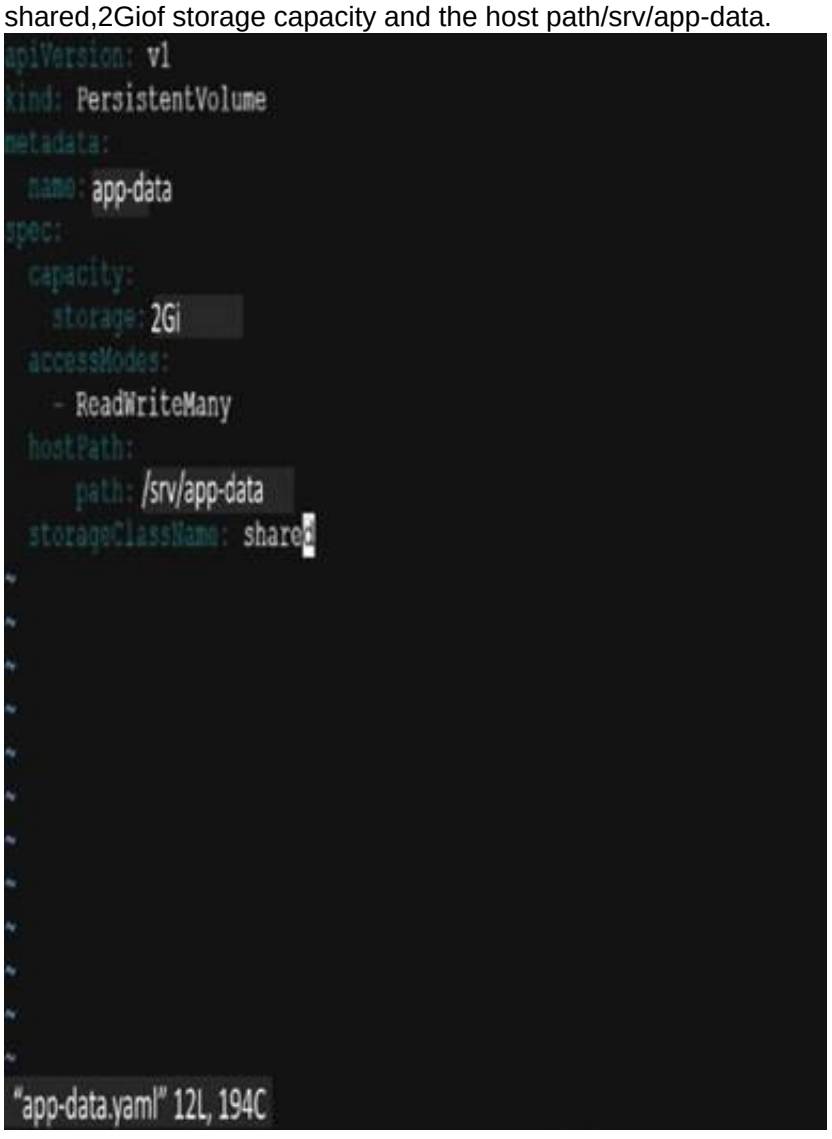

* 2. Save the file and create the persistent volume. Image for post

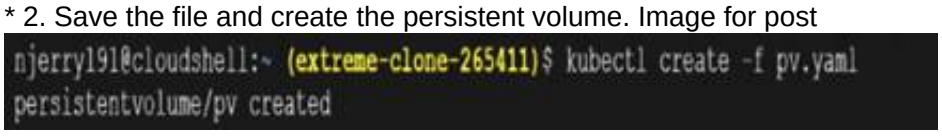

* 3. View the persistent volume.

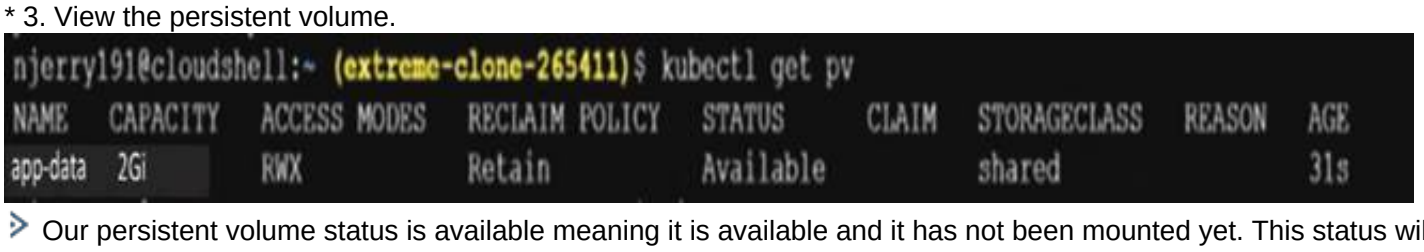

> Our persistent volume status is available meaning it is available and it has not been mounted yet. This status willchange when we mount the persistentVolume to a persistentVolumeClaim.
PersistentVolumeClaim
In a real ecosystem, a system admin will create the PersistentVolume then a developer will create a PersistentVolumeClaim which will be referenced in a pod. A PersistentVolumeClaim is created by specifying the minimum size and the access mode they require from the persistentVolume.
Challenge
> Create a Persistent Volume Claim that requests the Persistent Volume we had created above. The claim should request 2Gi. Ensurethat the Persistent Volume Claim has the same storageClassName as the persistentVolume you had previously created.
kind: PersistentVolumeapiVersion: v1metadata:name:app-data spec:
accessModes:-ReadWriteManyresources:

requests:storage:2Gi storageClassName:shared
* 2. Save and create the pvc
njerry191@cloudshell:~(extreme-clone-2654111)$ kubect1 create -f app-data.yaml persistentvolumeclaim/app-data created
* 3. View the pvc Image for post

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pvc
NAME    STATUS    VOLUME    CAPACITY    ACCESS MODES    STORAGECLASS
pv      Bound     pv        512m        RWX             shared
```

* 4. Let??s see what has changed in the pv we had initially created.
Image for post

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pv
NAME   CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM      STORAGECLASS   REASON   AGE
pv     512m       RWX            Retain           Bound    default/pv shared         16m
```

Our status has now changed fromavailabletobound.
* 5. Create a new pod named myapp with image nginx that will be used to Mount the Persistent Volume Claim with the path /var/app/config.
Mounting a Claim
apiVersion: v1kind: Podmetadata:creationTimestamp: nullname: app-dataspec:volumes:- name:congigpvcpersistenVolumeClaim:claimName: app-datacontainers:- image: nginxname: appvolumeMounts:- mountPath: "/srv/app-data"name: configpvc

NEW QUESTION 9
Create a namespace called 'development' and a pod with image nginx called nginx on this namespace.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl create namespace development
kubectl run nginx --image=nginx --restart=Never -n development

NEW QUESTION 10
Create a pod as follows:
- Name:non-persistent-redis
- container Image:redis
- Volume with name:cache-control
- Mount path:/data/redis

The pod should launch in thestagingnamespace and the volumemust notbe persistent.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\13 B.JPG

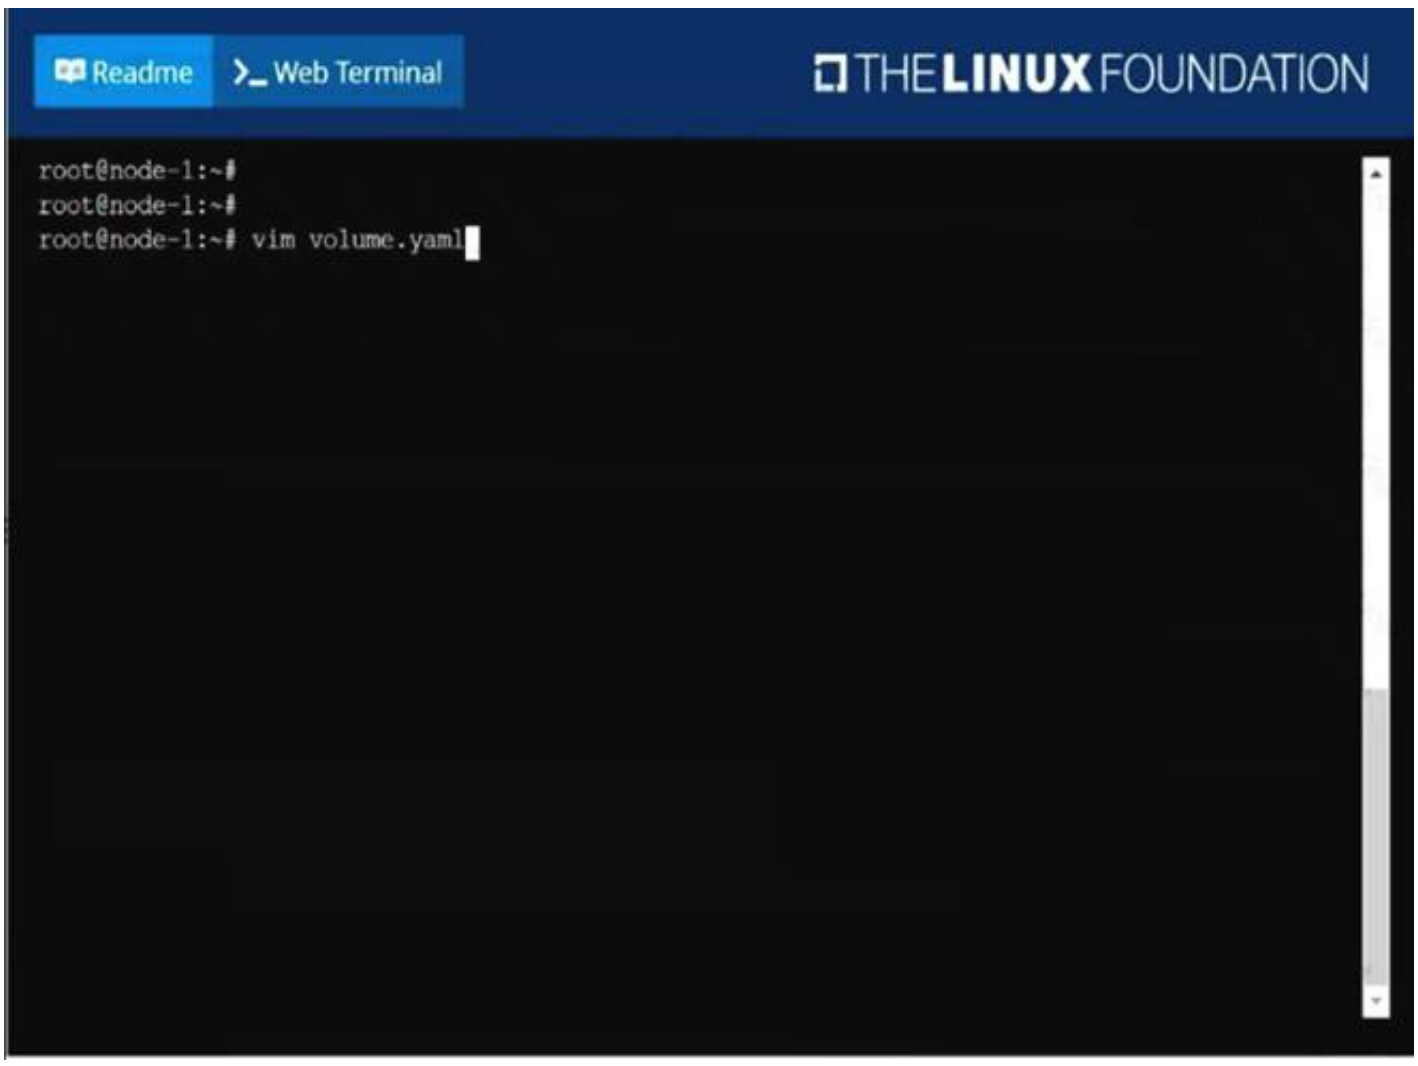

F:\Work\Data Entry Work\Data Entry\20200827\CKA\13 C.JPG

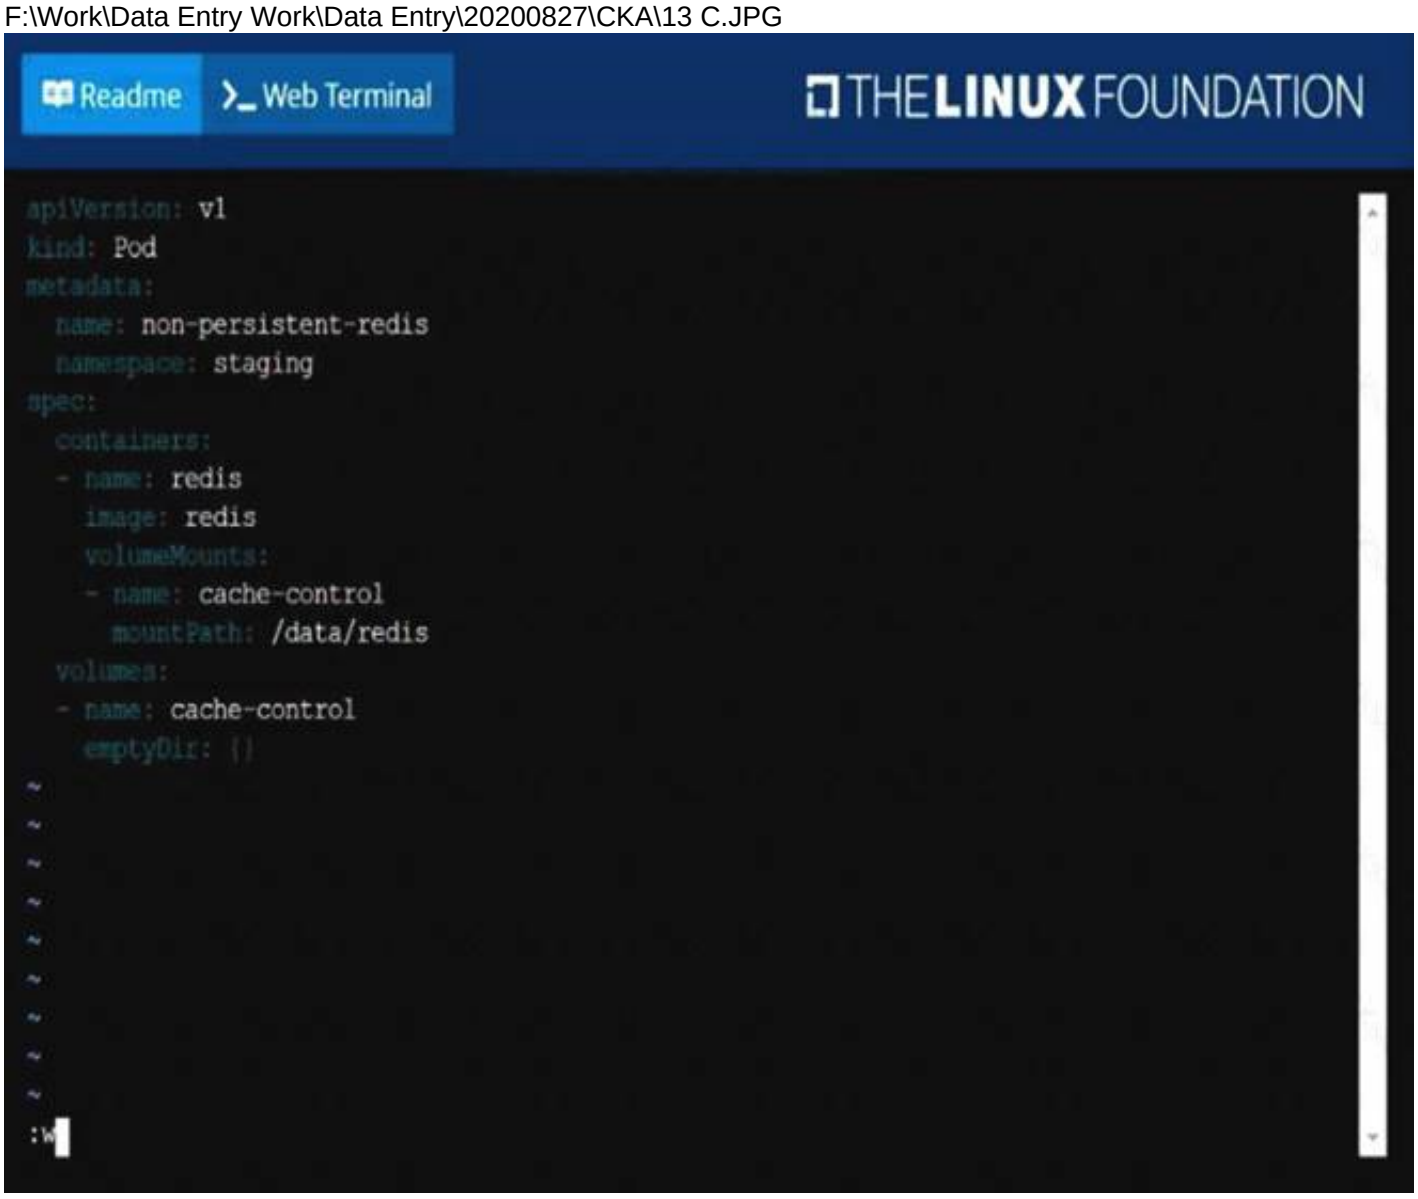

F:\Work\Data Entry Work\Data Entry\20200827\CKA\13 D.JPG

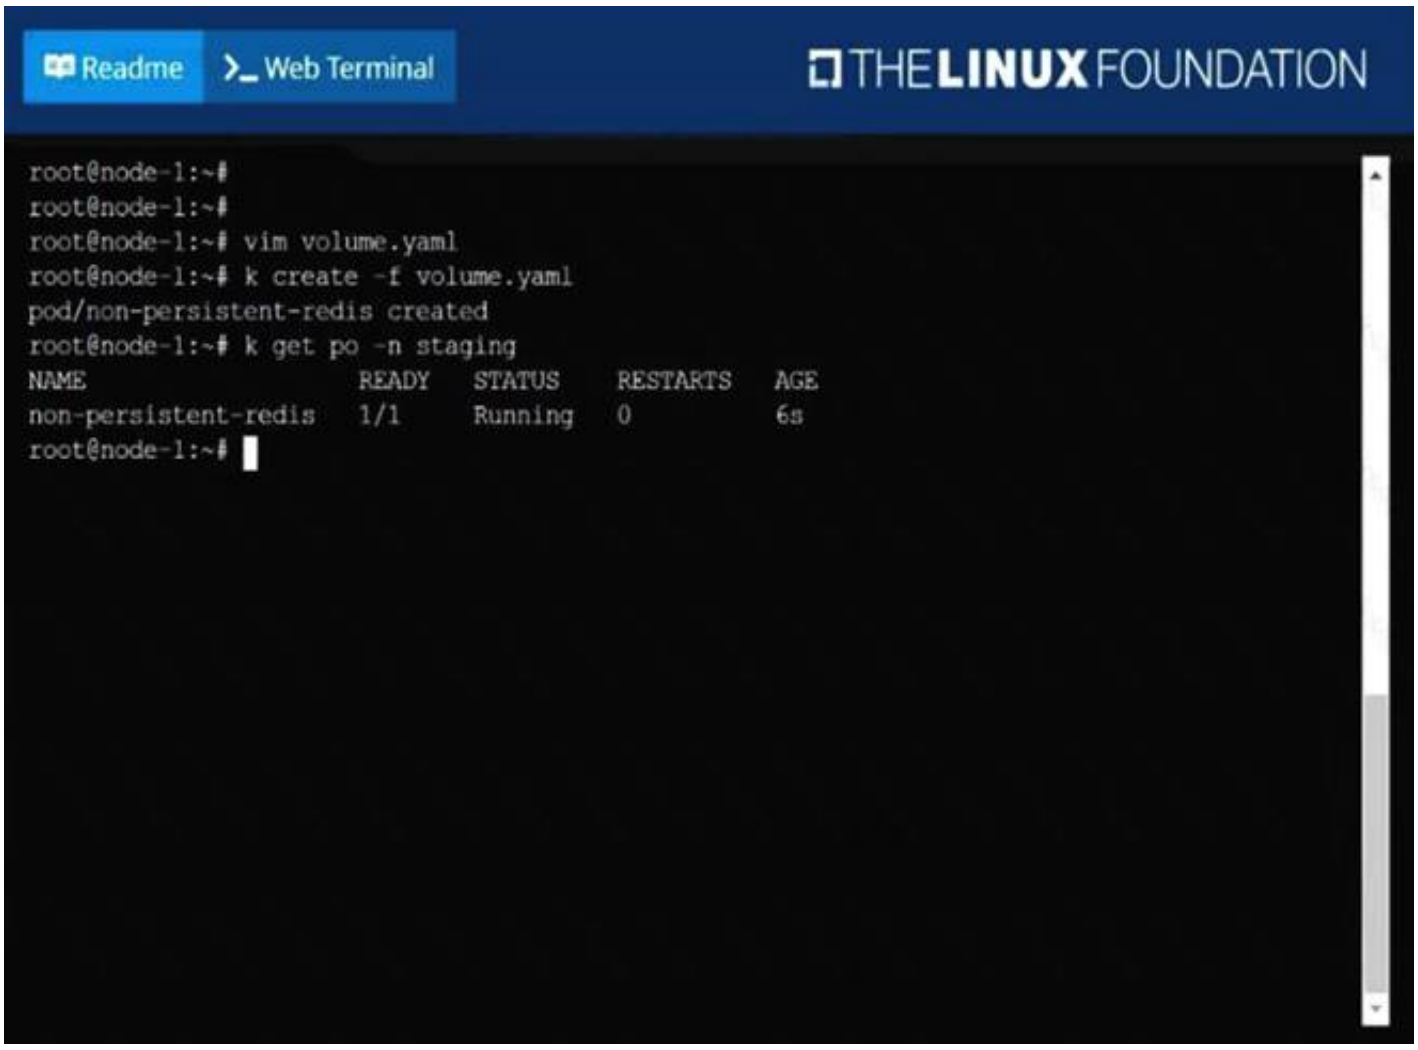

NEW QUESTION 10
List the nginx pod with custom columns POD_NAME and POD_STATUS

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl get po -o=custom-columns="POD_NAME:.metadata.name, POD_STATUS:.status.containerStatuses[].state"

NEW QUESTION 12
Ensure a single instance of podnginxis running on each node of theKubernetes cluster wherenginxalso represents the Image name whichhas to be used. Do not override anytaints currently in place.
UseDaemonSetto complete thistask and useds-kusc00201asDaemonSet name.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\3 B.JPG

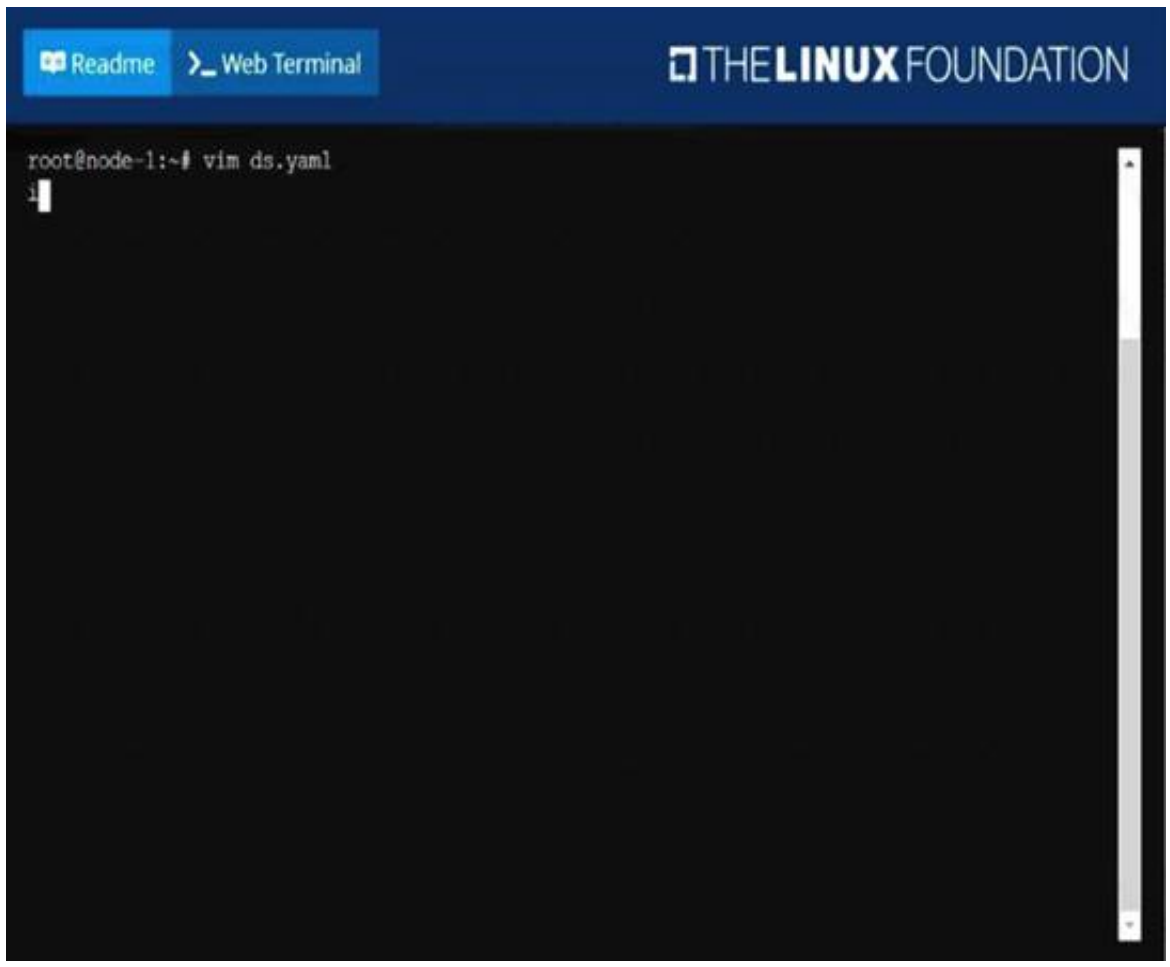

F:\Work\Data Entry Work\Data Entry\20200827\CKA\3 C.JPG

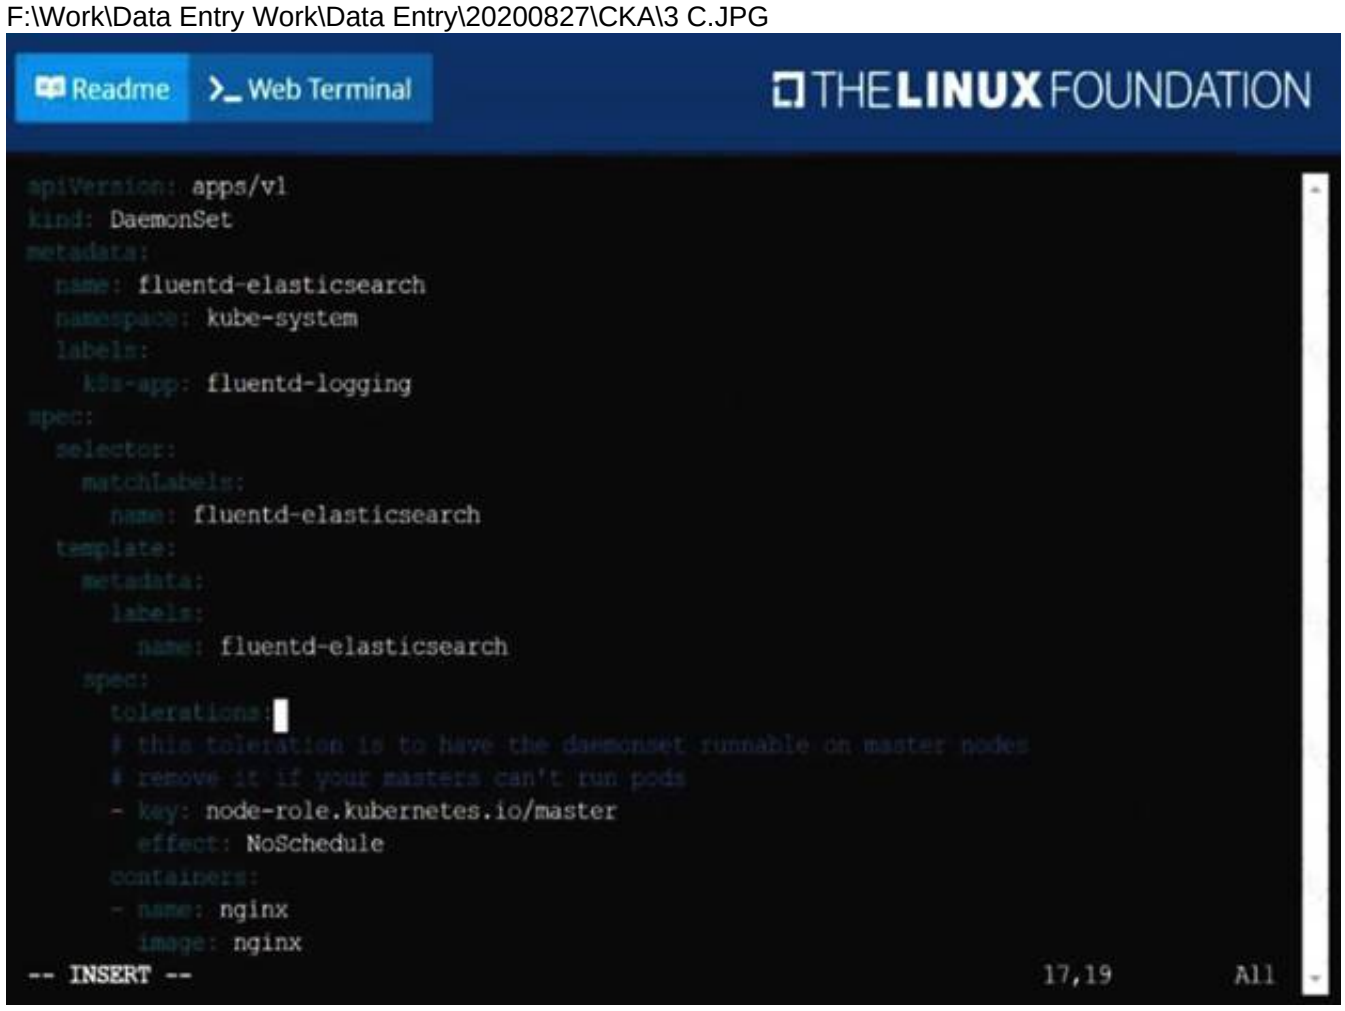

F:\Work\Data Entry Work\Data Entry\20200827\CKA\3 D.JPG

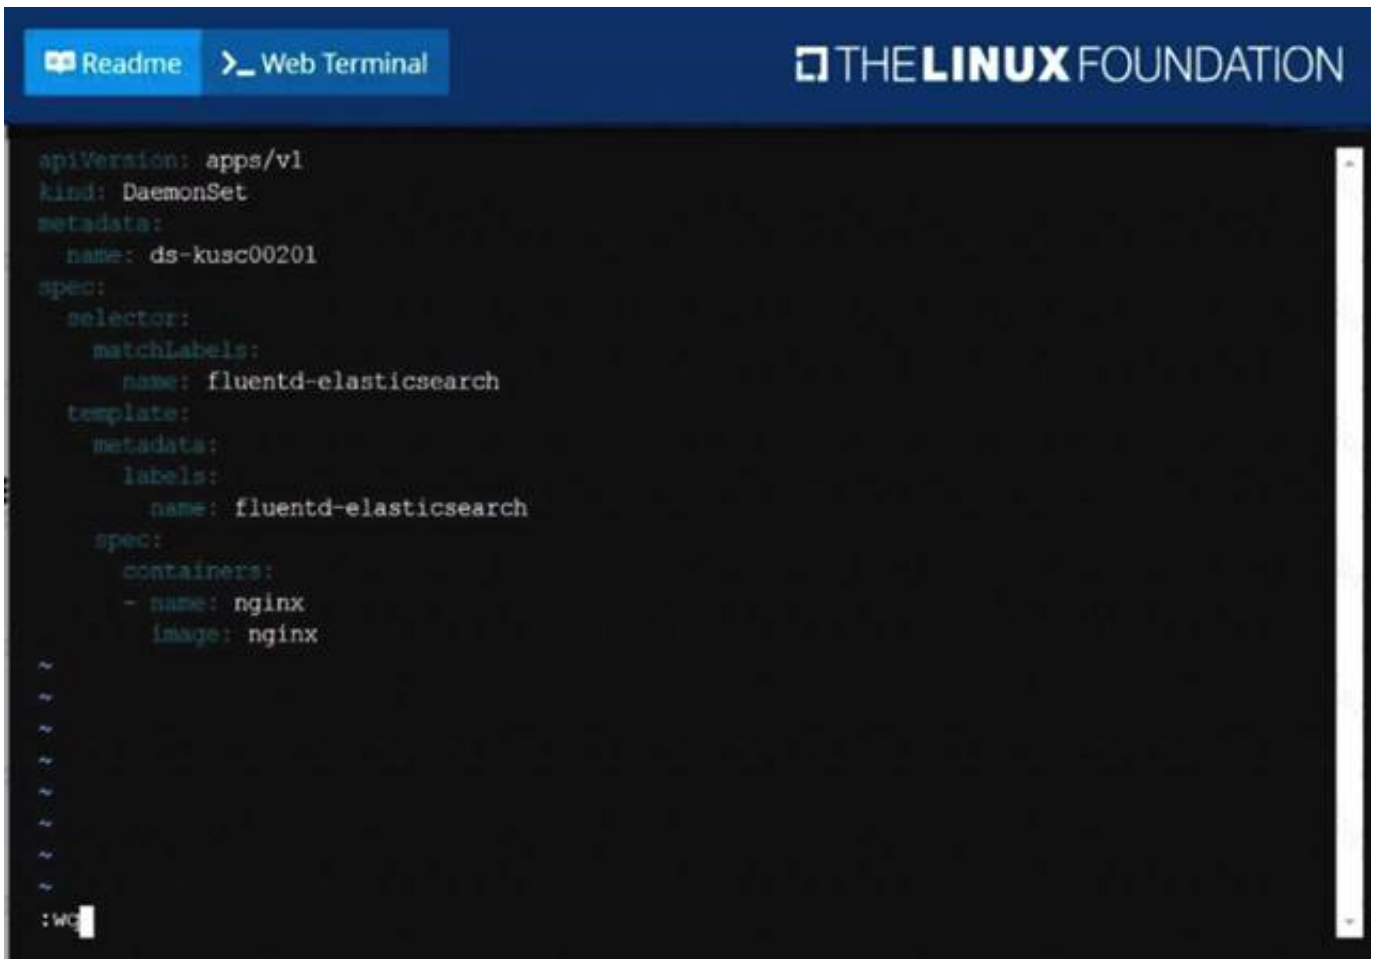

F:\Work\Data Entry Work\Data Entry\20200827\CKA\3 E.JPG

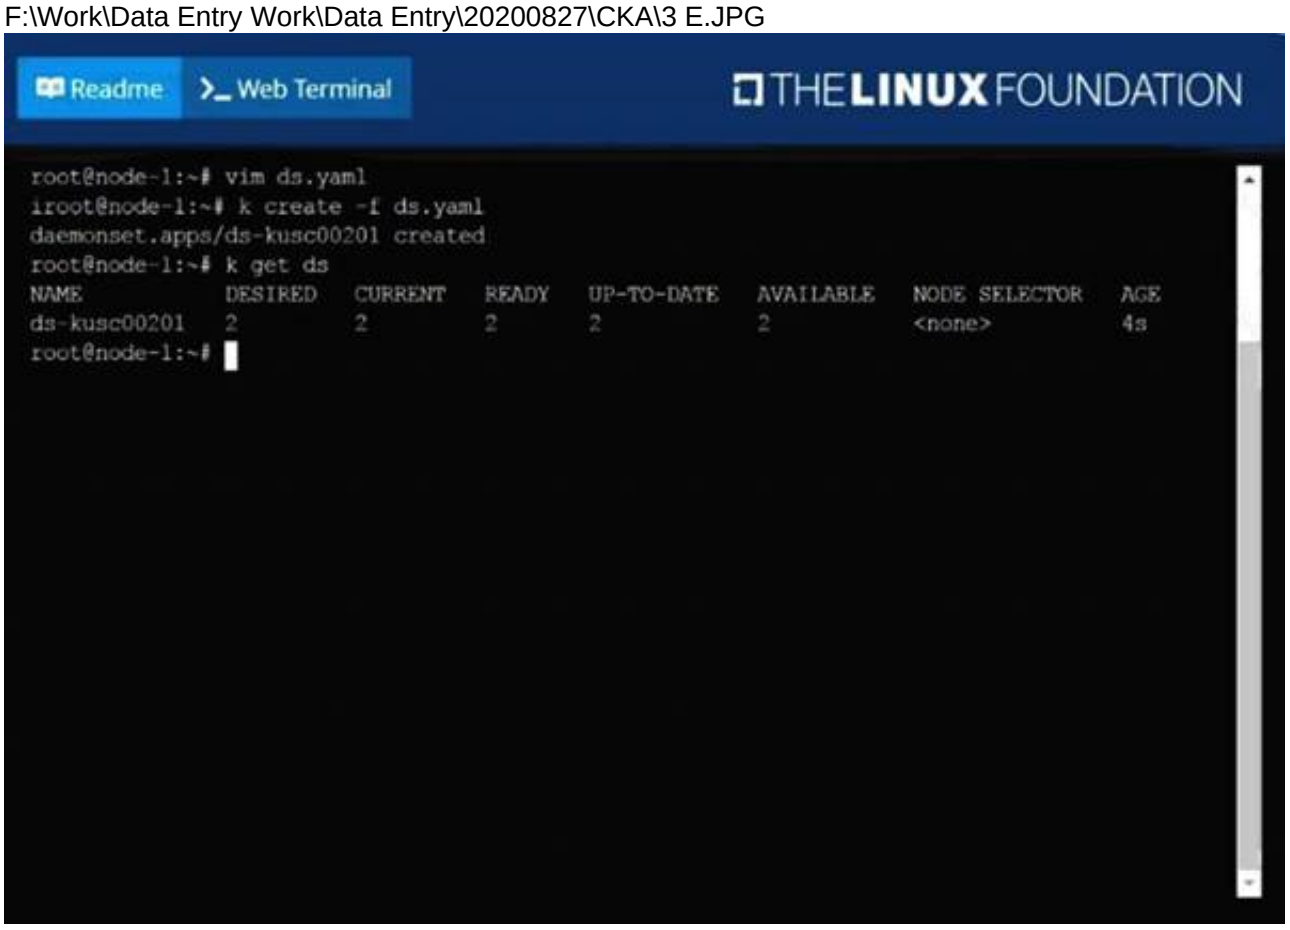

NEW QUESTION 14
Check the image version in pod without the describe command

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl get po nginx -o jsonpath='{.spec.containers[].image}{"\n"}'

NEW QUESTION 16
Configure the kubelet systemd-managed service, on the nodelabelled withname=wk8s-node-1, tolaunch a pod containing a singlecontainer of Imagehttpdnamedwebtoolautomatically. Any spec filesrequired should be placed in the/etc/kubernetes/manifestsdirectoryon the node.
You canssh to theappropriate node using:
[student@node-1] $ sshwk8s-node-1
You can assume elevatedprivileges on the node with thefollowing command:
[student@wk8s-node-1] $ |sudo ?Ci

A. Mastered
B. Not Mastered

Answer: A

Explanation:

solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\21 C.JPG

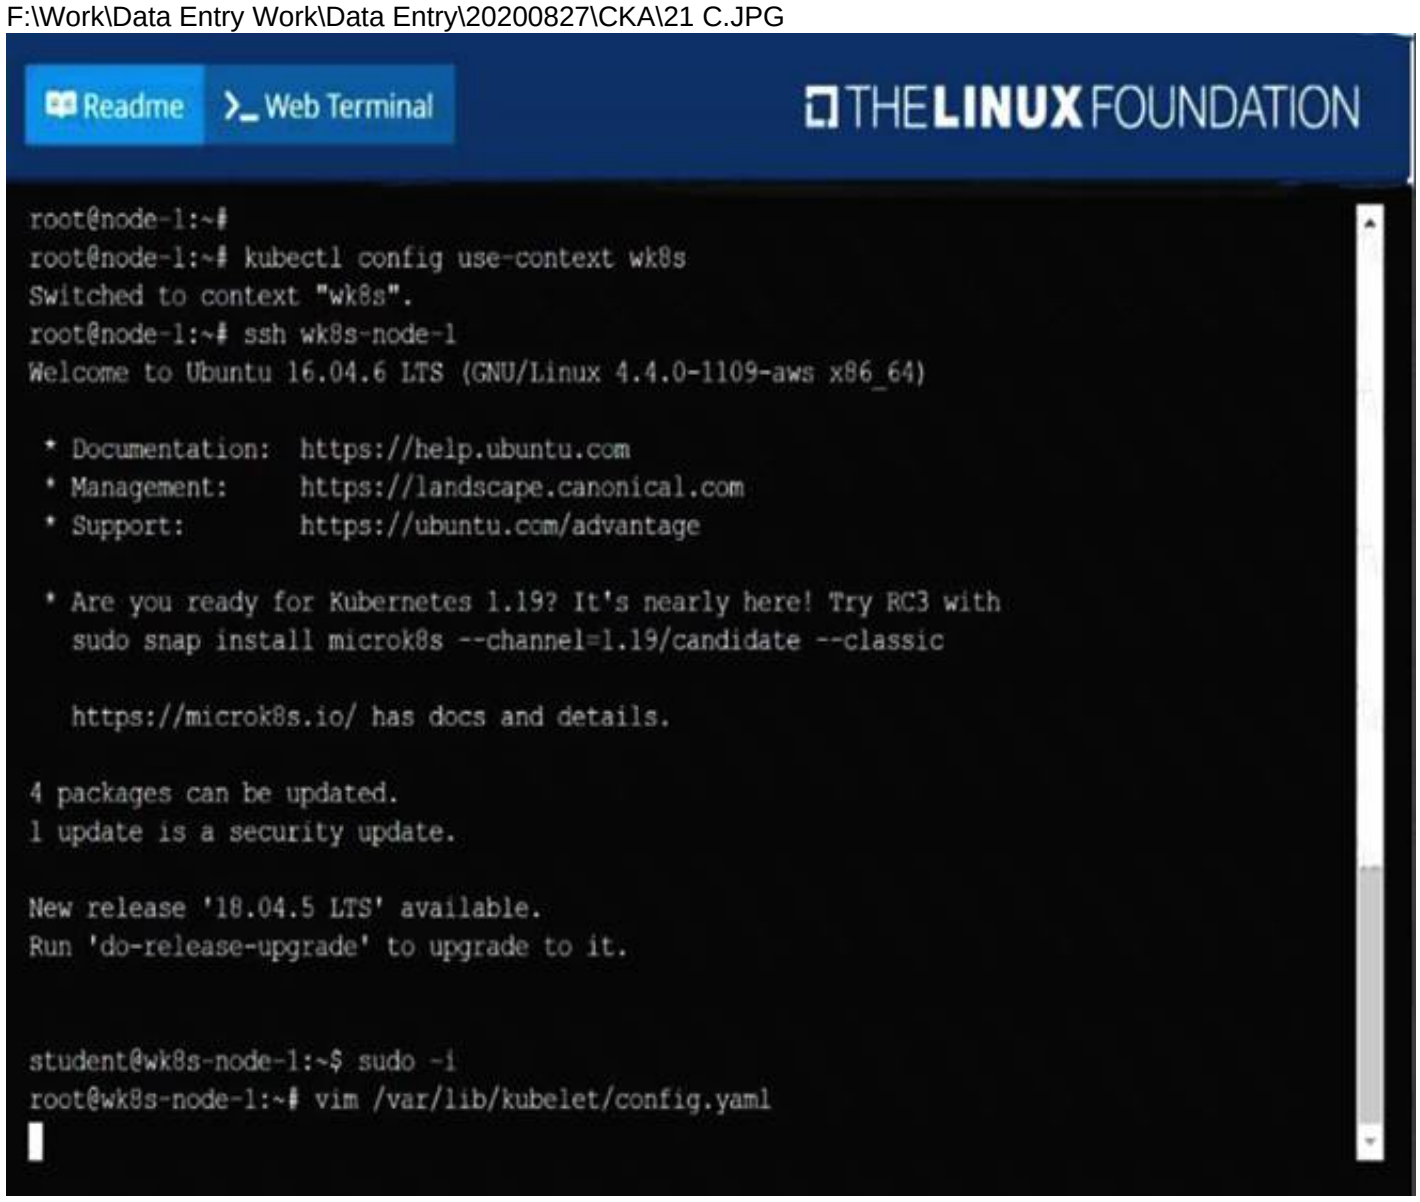

F:\Work\Data Entry Work\Data Entry\20200827\CKA\21 D.JPG

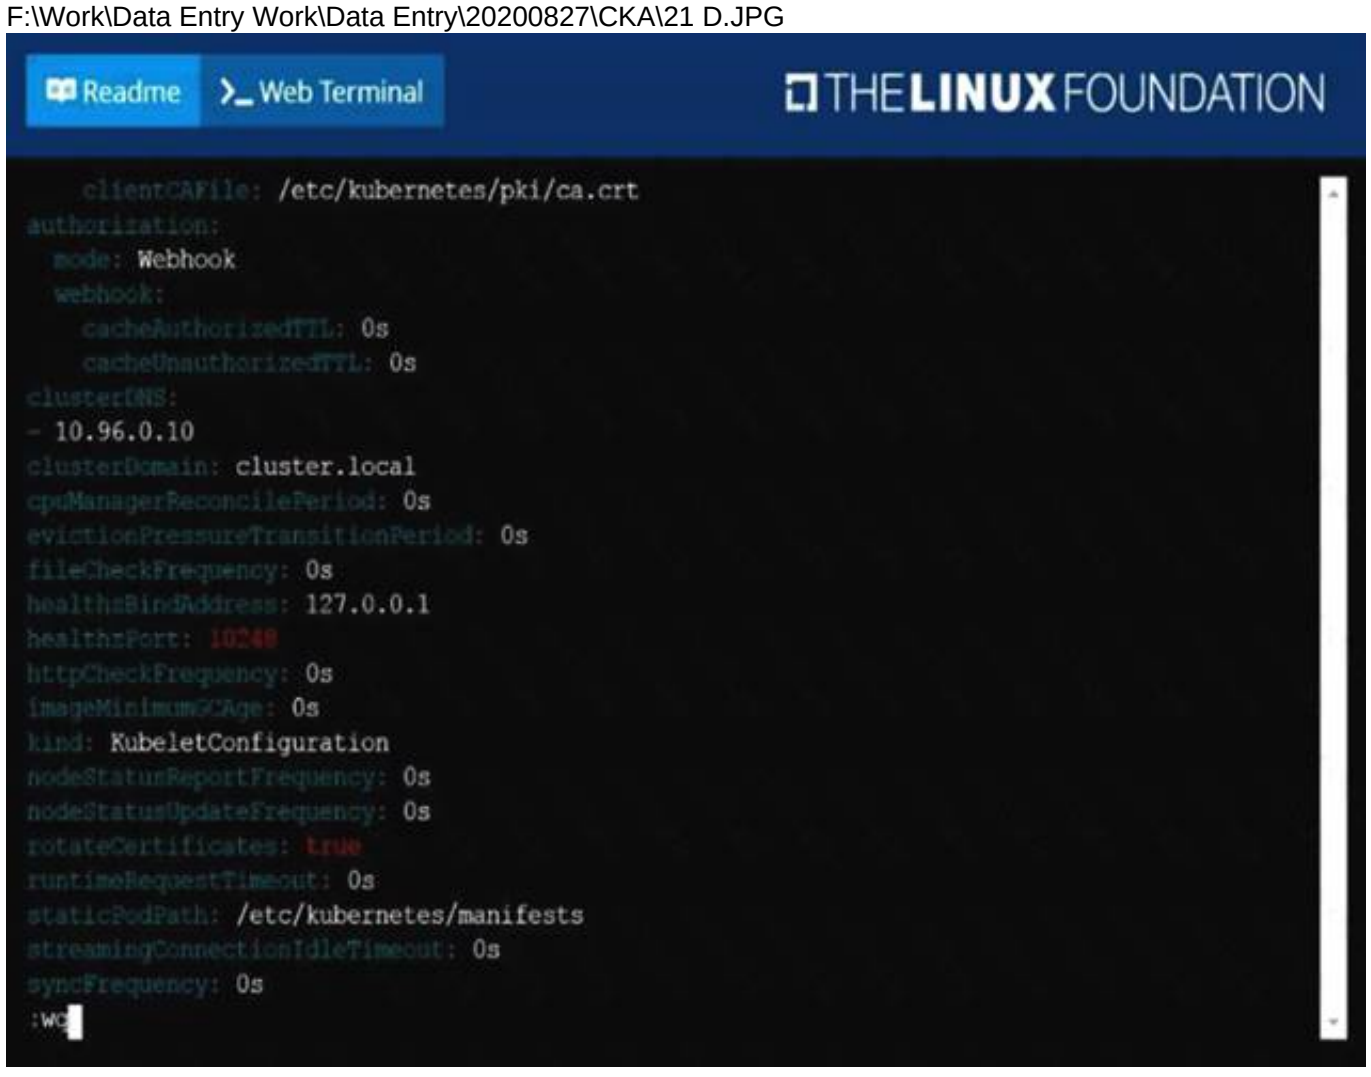

F:\Work\Data Entry Work\Data Entry\20200827\CKA\21 E.JPG

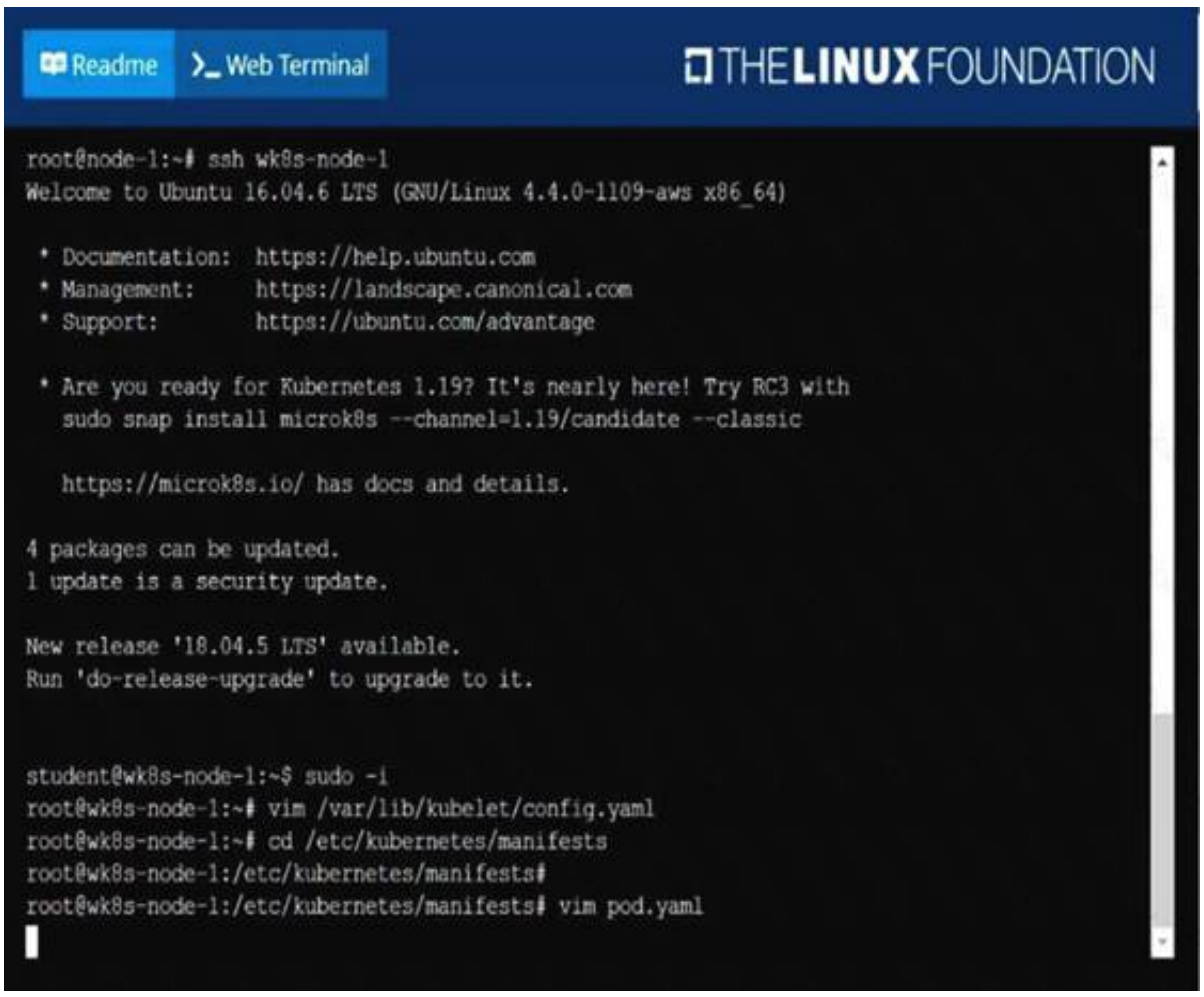

F:\Work\Data Entry Work\Data Entry\20200827\CKA\21 F.JPG

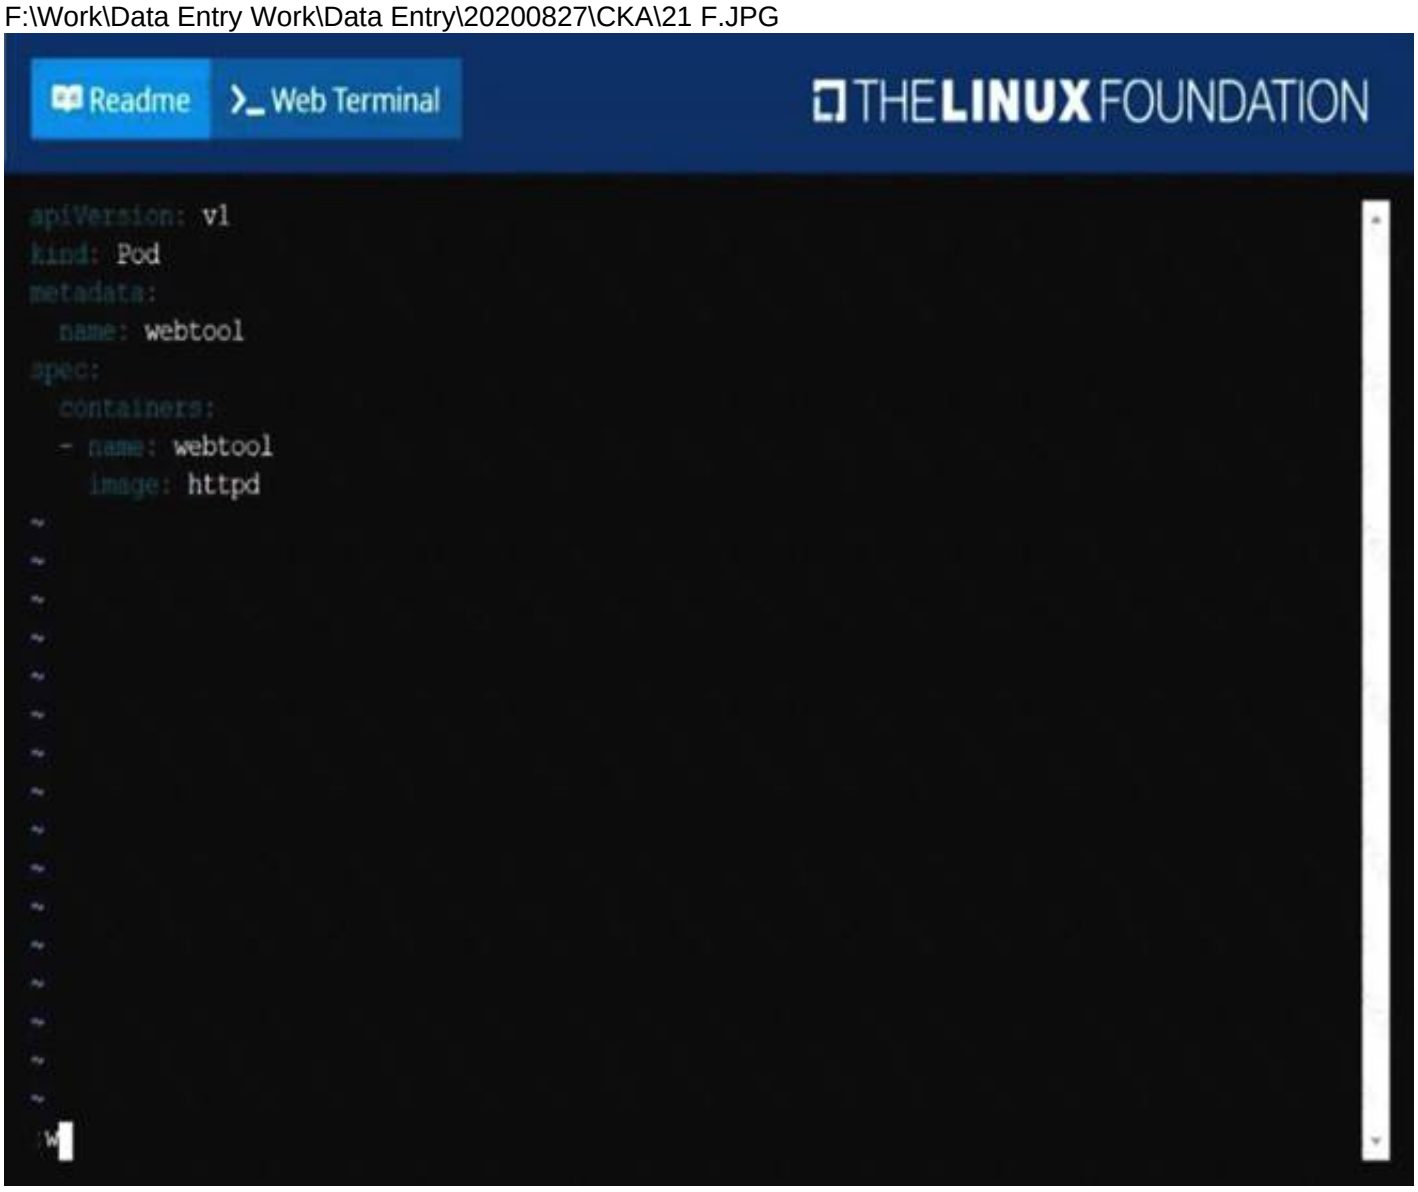

F:\Work\Data Entry Work\Data Entry\20200827\CKA\21 G.JPG

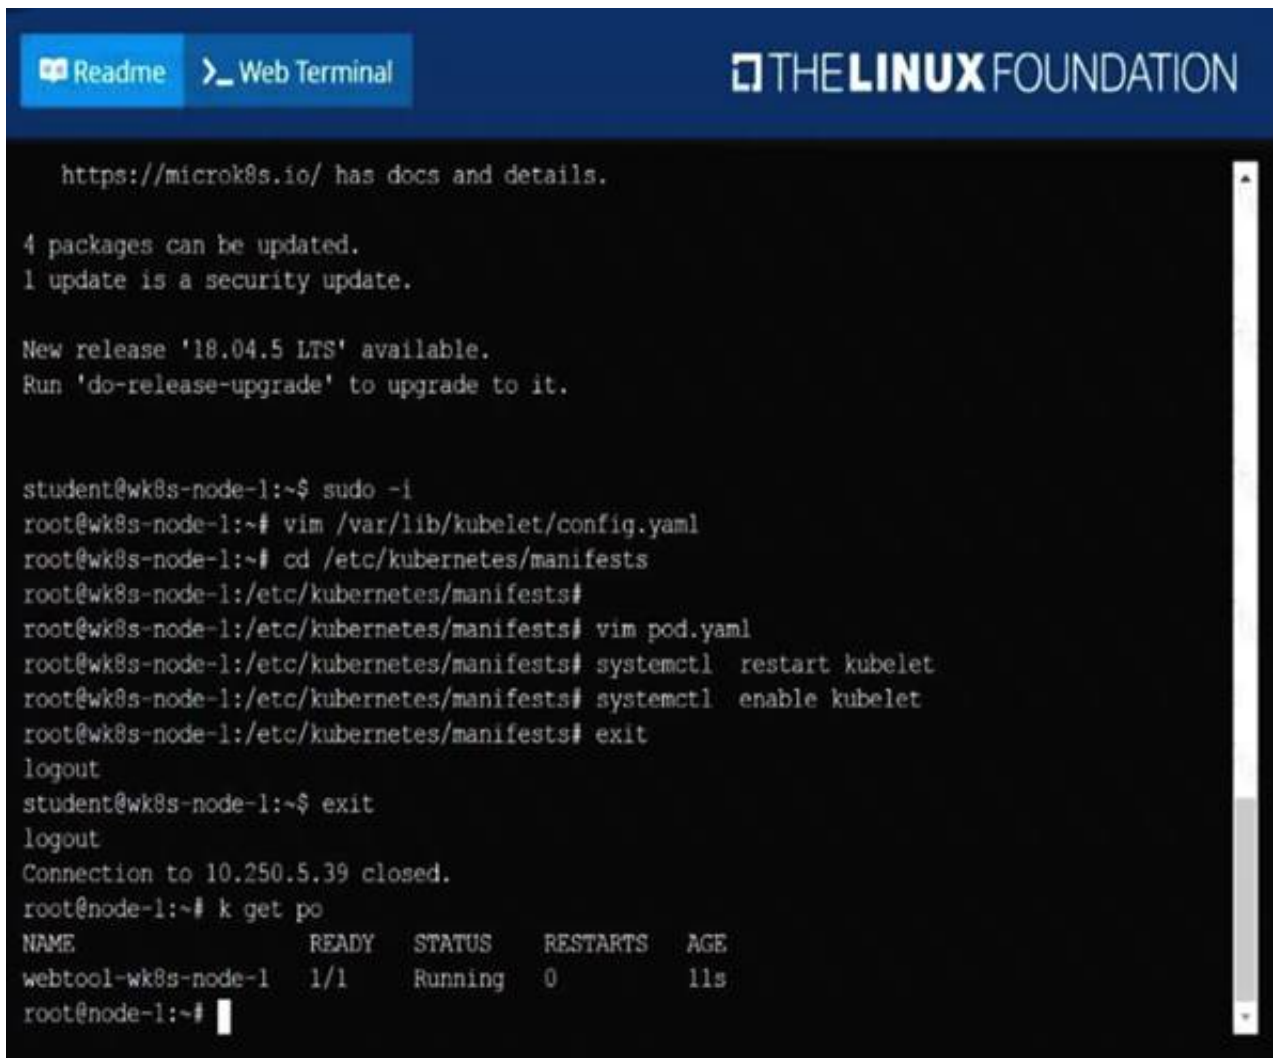

NEW QUESTION 18
Perform the following tasks:
- Add an init container tohungry-bear(which has beendefined in spec file /opt/KUCC00108/pod-spec-KUCC00108.yaml)
- The init container should createan empty file named/workdir/calm.txt
- If/workdir/calm.txtis notdetected, the pod should exit
- Once the spec file has beenupdatedwith the init containerdefinition, the pod should becreated

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\4 B.JPG

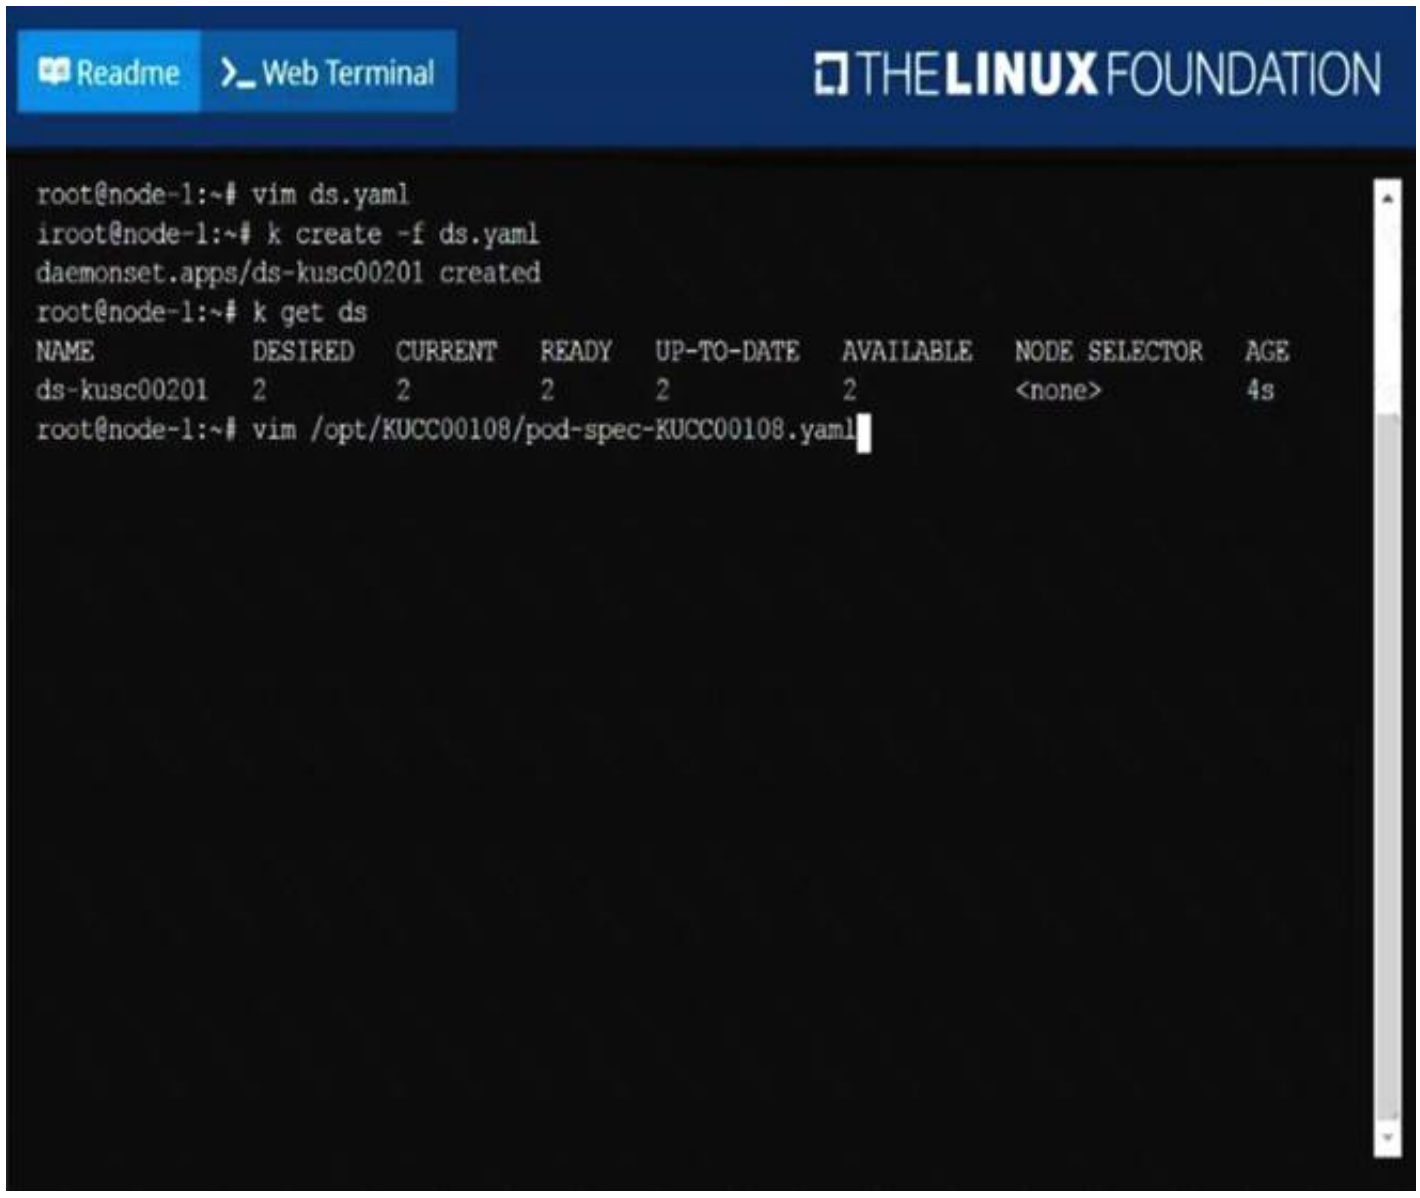

F:\Work\Data Entry Work\Data Entry\20200827\CKA\4 C.JPG

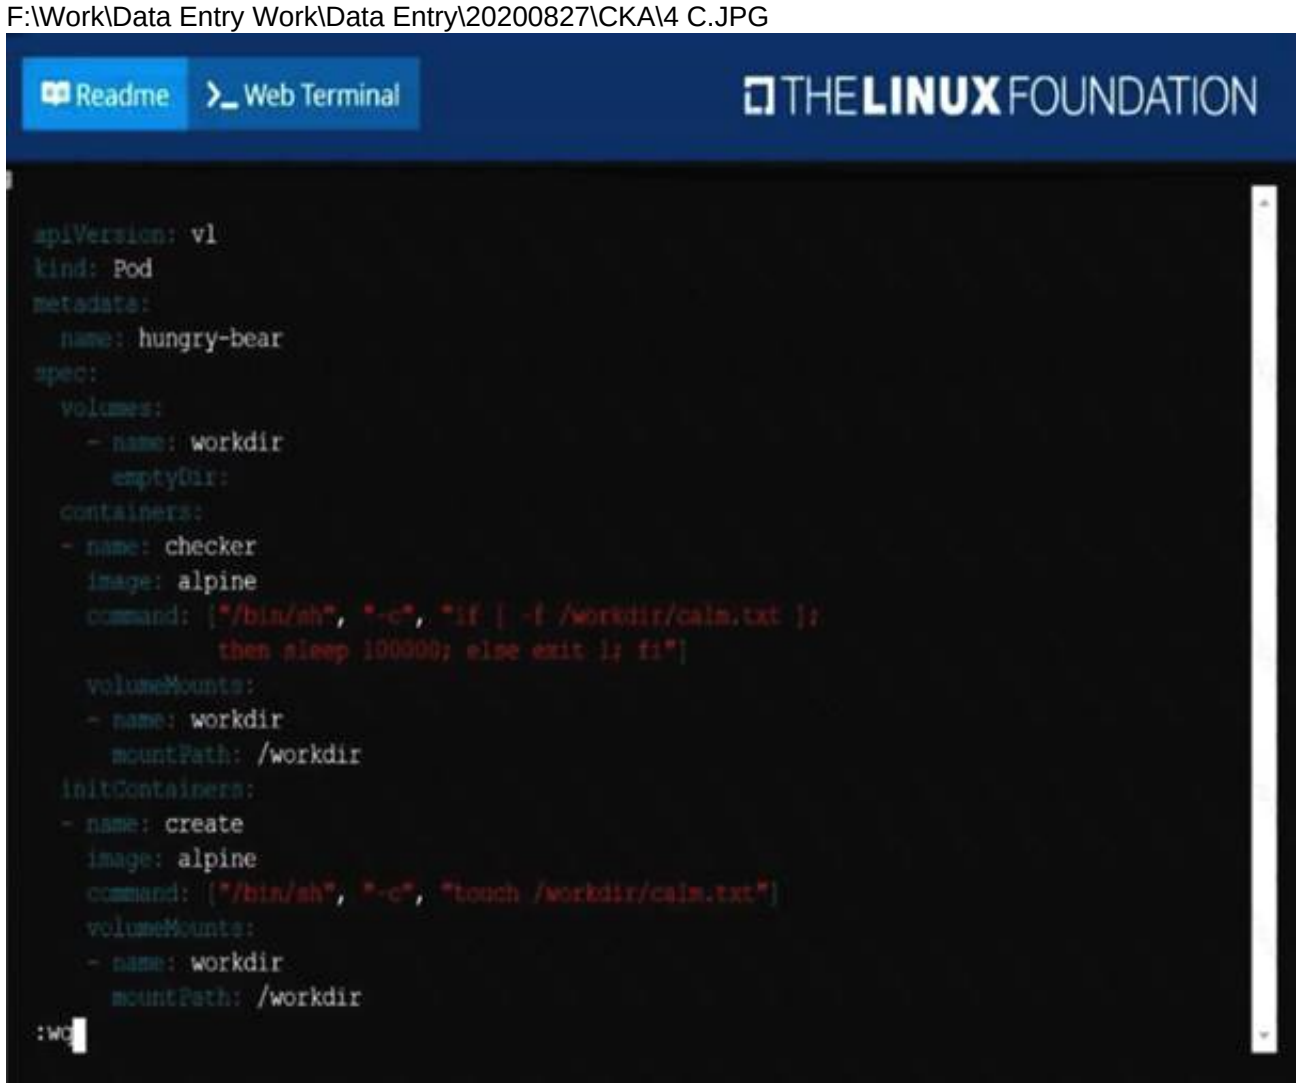

F:\Work\Data Entry Work\Data Entry\20200827\CKA\4 D.JPG

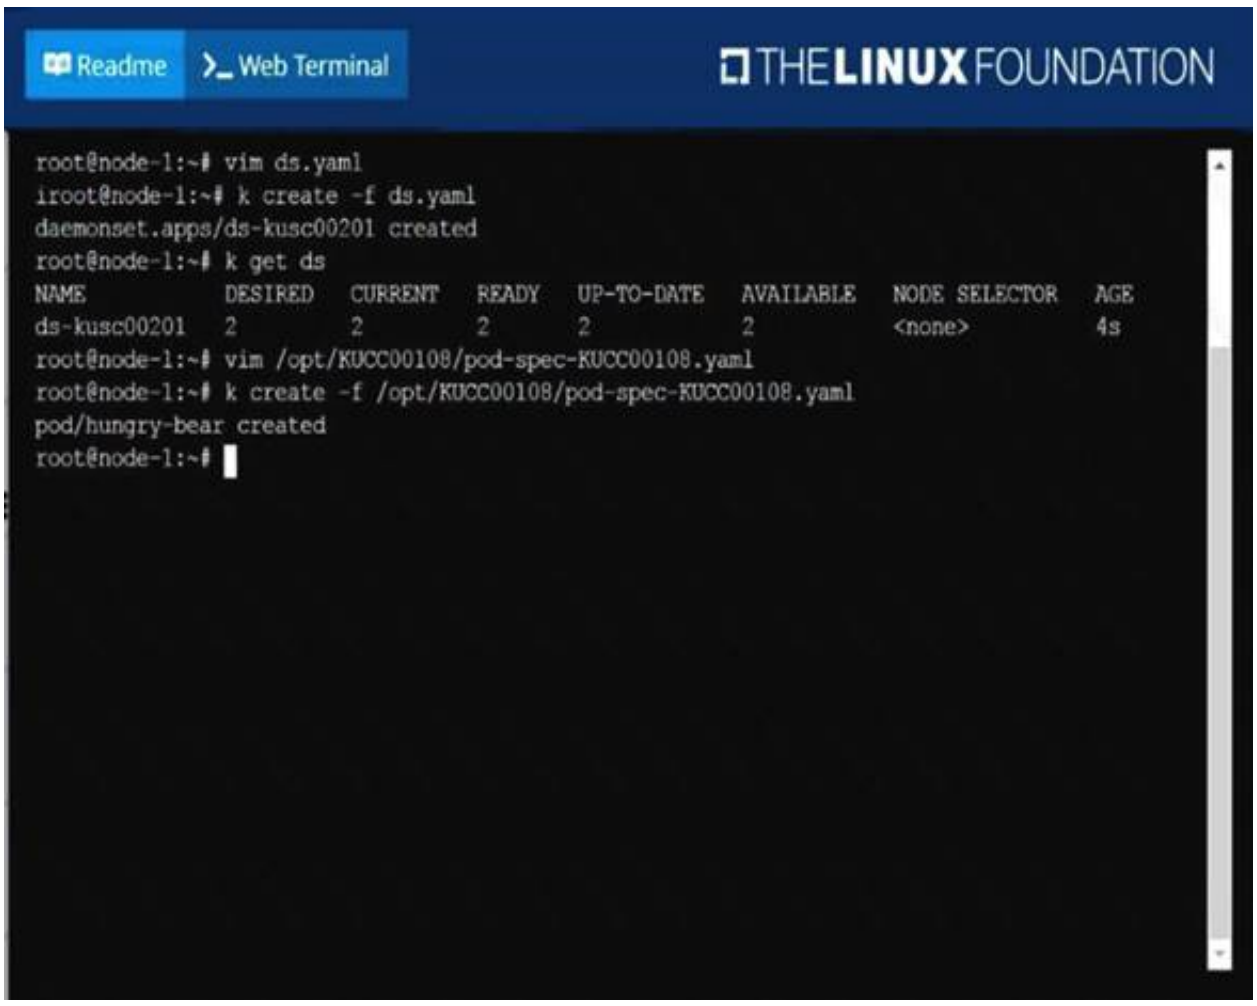

NEW QUESTION 21
Create a busybox pod and add ??sleep 3600?? command

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl run busybox --image=busybox --restart=Never -- /bin/sh -c "sleep 3600"

NEW QUESTION 26
Create and configure the servicefront-end-serviceso it's accessiblethroughNodePortand routes to theexisting pod namedfront-end.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\8 B.JPG

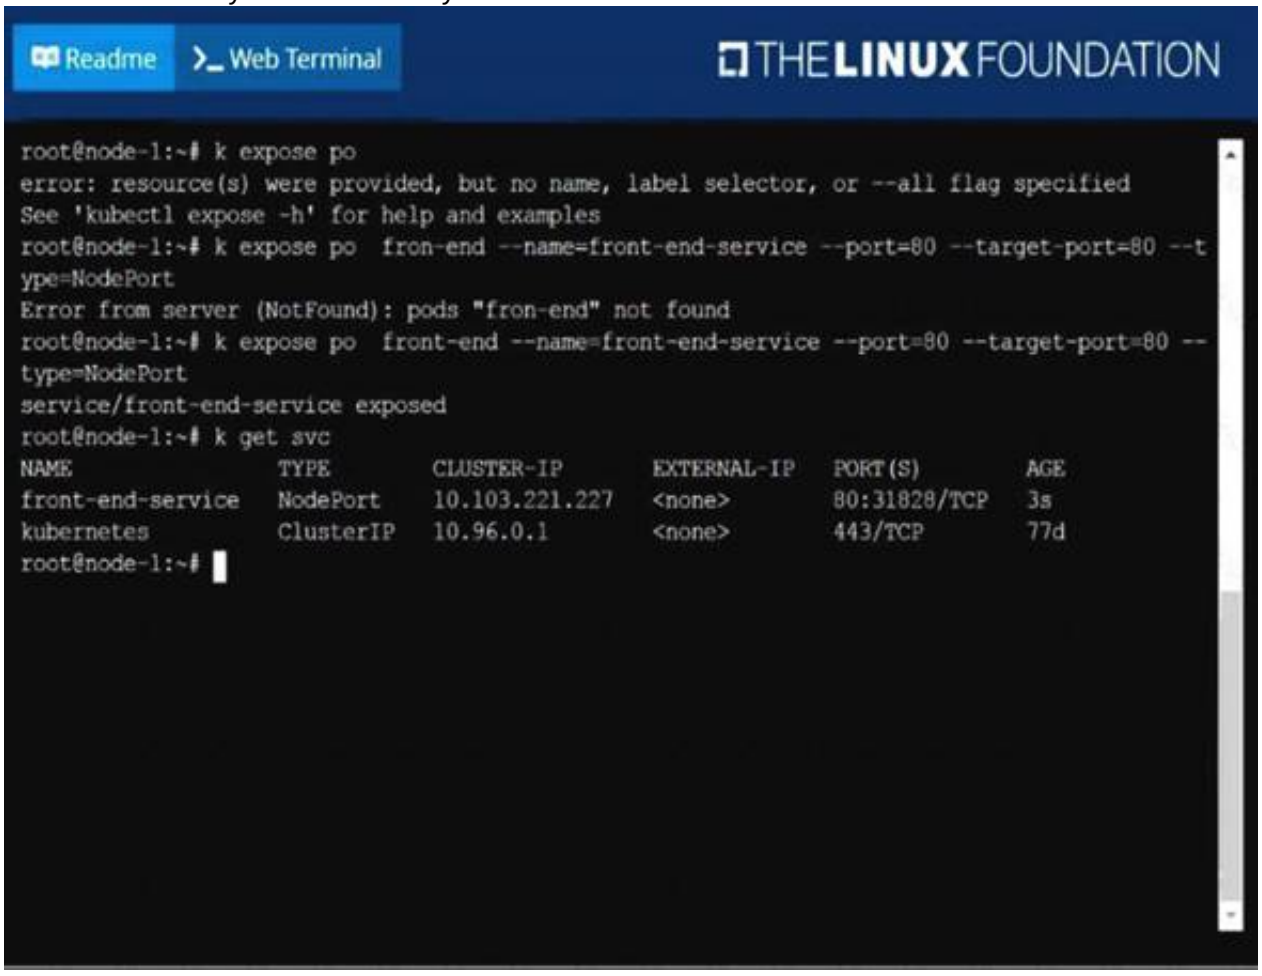

NEW QUESTION 29
List all the pods sorted by name

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubectl get pods --sort-by=.metadata.name

NEW QUESTION 33
For this item, you will havetosshto the nodesik8s-master-0andik8s-node-0and complete all tasks on thesenodes. Ensure that you return tothe base node (hostname:node-1) when you havecompleted this item.
Context
As an administrator of a smalldevelopment team, you have beenasked to set up a Kubernetes clusterto test the viability of a newapplication.
Task
You must usekubeadmto performthis task. Anykubeadminvocationswill require the use of the
--ignore-preflight-errors=alloption.
- Configure thenodeik8s-master-Oas a masternode. .
- Join the nodeik8s-node-otothe cluster.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
You must use thekubeadmconfiguration file located at/etc/kubeadm.confwhen initializingyour cluster.
You may use any CNI pluginto complete this task, but ifyou don't have your favouriteCNI plugin's manifest URL athand, Calico is one popularoption:https://docs.projectcalico.org/v3.14/manifests/calico.yaml
Docker is already installedon both nodes andapthasbeen configured so that you caninstall the required tools.

NEW QUESTION 34
Get IP address of the pod ?C ??nginx-dev??

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Kubect1 get po -o wide Using JsonPath
kubect1 get pods -o=jsonpath='{range items[*]}{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'

NEW QUESTION 36
Scale the deploymentwebserverto6pods.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\14 B.JPG

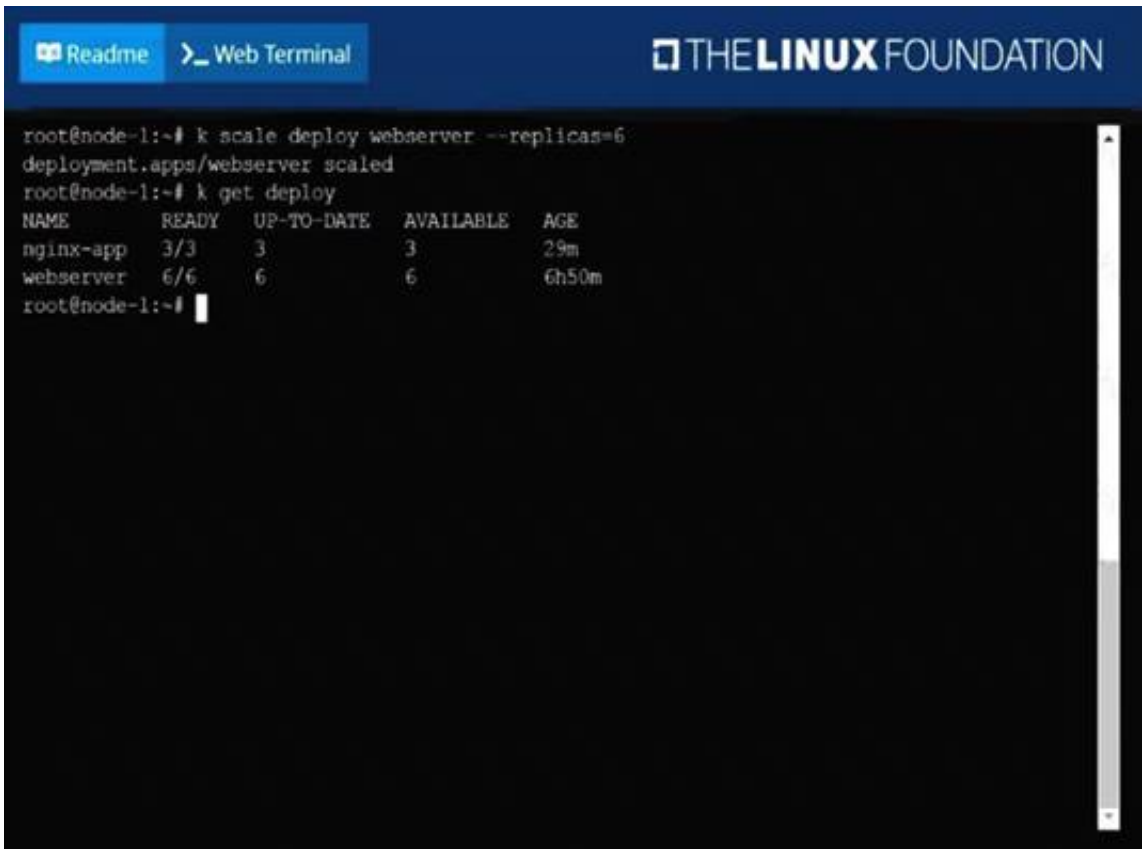

NEW QUESTION 39
List all the pods sorted by created timestamp

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubect1 get pods--sort-by=.metadata.creationTimestamp

NEW QUESTION 43
Print pod name and start time to ??/opt/pod-status?? file

A. Mastered
B. Not Mastered

Answer: A

Explanation:
kubect1 get pods -o=jsonpath='{range items[*]}{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'

NEW QUESTION 47
Check to see how many worker nodes are ready (not including nodes taintedNoSchedule) and write the number to/opt/KUCC00104/kucc00104.txt.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
solution
F:\Work\Data Entry Work\Data Entry\20200827\CKA\15 B.JPG

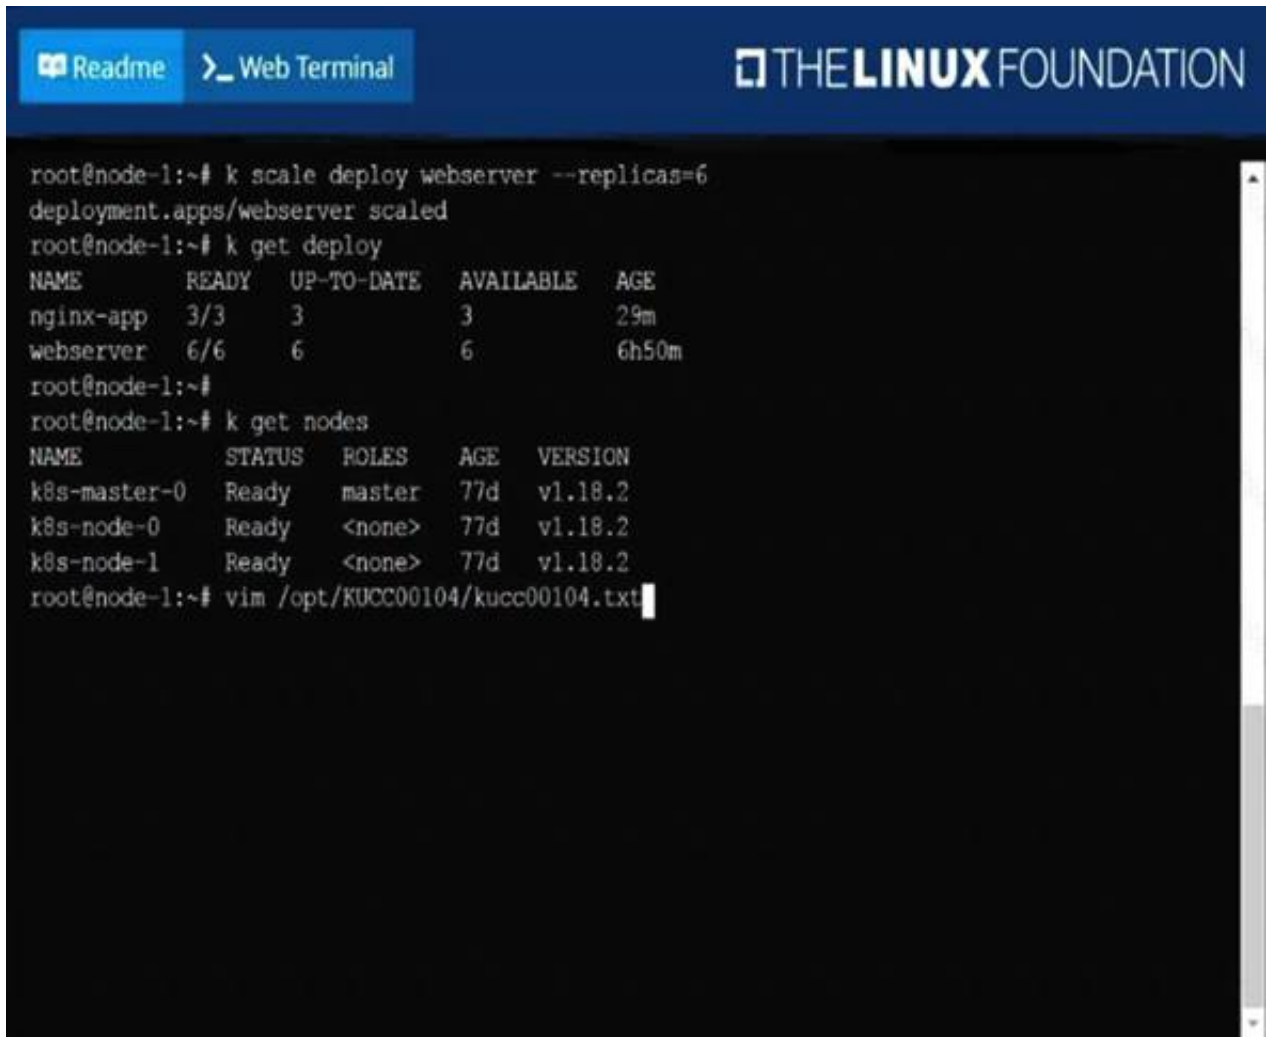

F:\Work\Data Entry Work\Data Entry\20200827\CKA\15 C.JPG

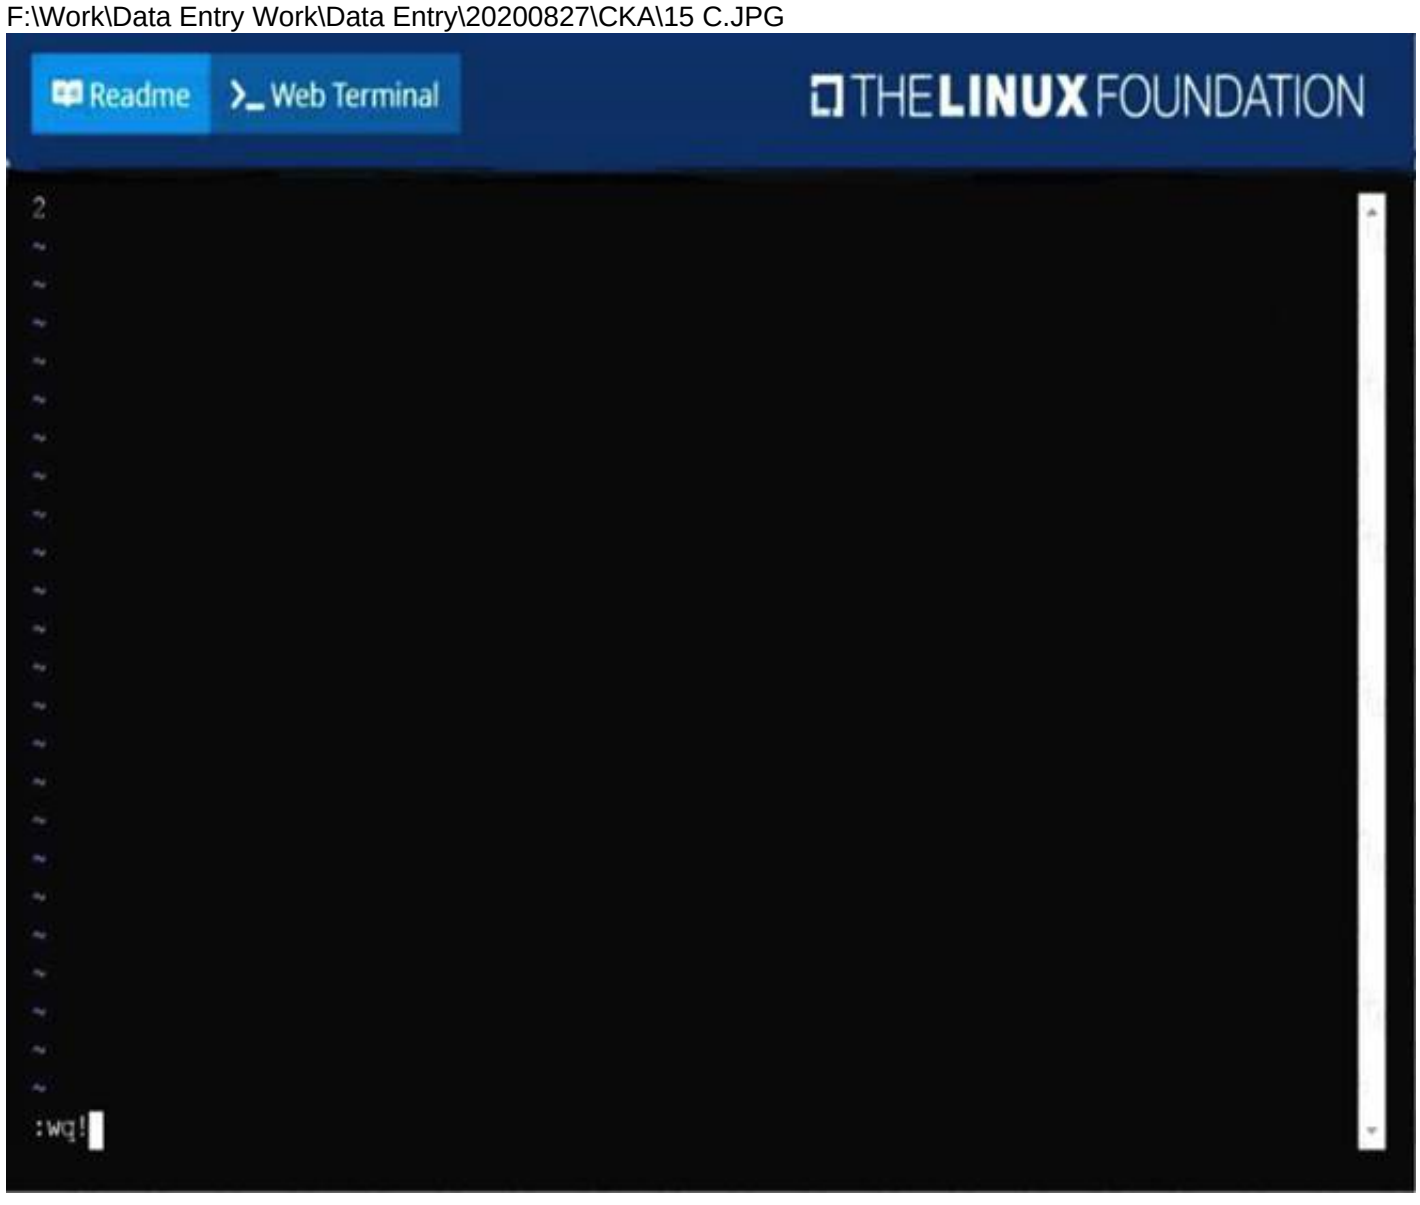

NEW QUESTION 51
.....

THANKS FOR TRYING THE DEMO OF OUR PRODUCT

Visit Our Site to Purchase the Full Set of Actual CKA Exam Questions With Answers.

We Also Provide Practice Exam Software That Simulates Real Exam Environment And Has Many Self-Assessment Features. Order the CKA Product From:

https://www.2passeasy.com/dumps/CKA/

Money Back Guarantee

CKA Practice Exam Features:

* CKA Questions and Answers Updated Frequently

* CKA Practice Questions Verified by Expert Senior Certified Staff

* CKA Most Realistic Questions that Guarantee you a Pass on Your FirstTry

* CKA Practice Test Questions in Multiple Choice Formats and Updatesfor 1 Year